





Kubernetes 101

Kubernetes 101



Damian Flynn

MVP Cloud & Datacenter + Cisco Champion



@Damian_Flynn

www.DamianFlynn.Com

NIC

Introductions

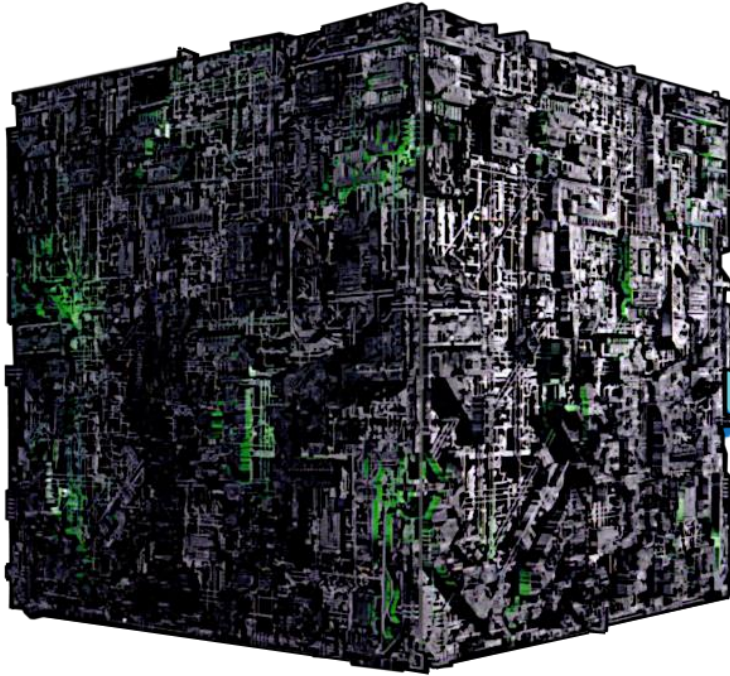
- Where did it come from
- What does the name mean
- What does it do
- Why do we have it

Introductions



- Born in Google
- 2014, donated to the CNCF
- Composed in Go/Golang
- Online Community
 - Git: <https://github.com/kubernetes/kubernetes>
 - Twitter: @kubernetesio
 - Slack: slack.k8s.io

Introductions



Borg
Proprietary



Omega
Proprietary



Kubernetes
Open Source



Kubernetes

- Greek for “Helmsman”
 - The person responsible for steering the ship

Introductions



~~Kubernetes~~

K8s

Introductions

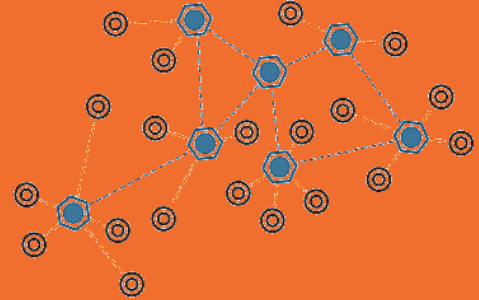
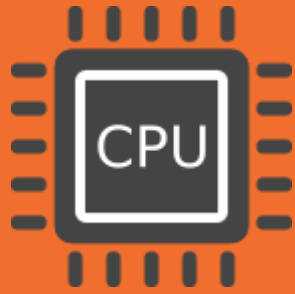


- Opensource (Apache 2.0 License)
- Shortened to 'K8s'
- Based on Googles Internal tools, Borg and Omega
- Donated to CNCF in 2014
- V1.0 release in July 2015
- Current release V1.8

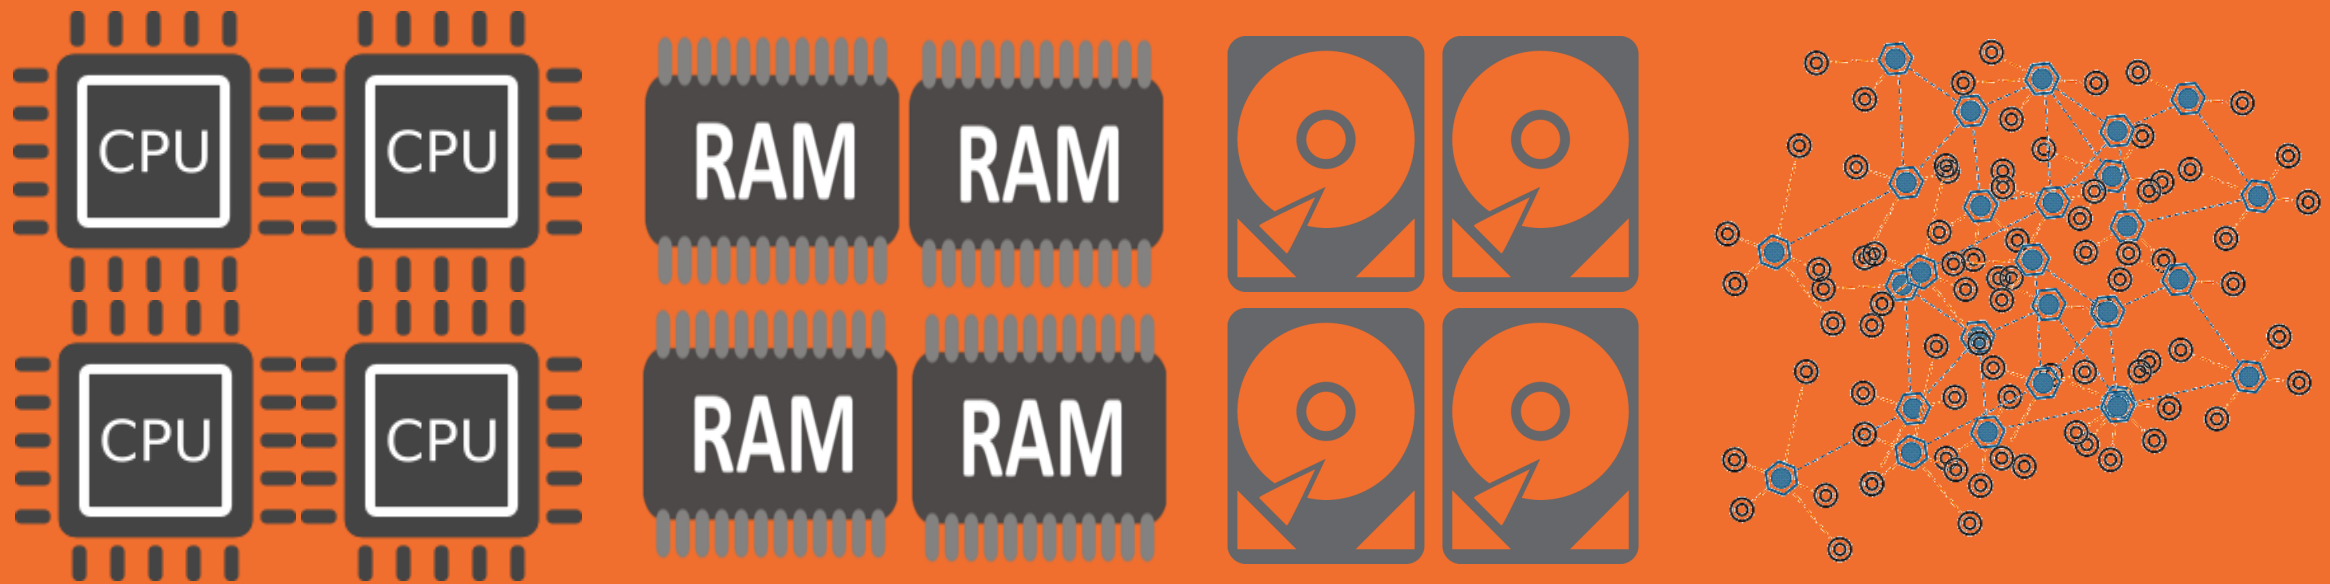
Containers

Scalability Challenges





Viewing the Data Center as a Computer



Viewing the Data Center as a Computer

Kubernetes can manage it

What is it?



- Standard Package format
- Packaging Manifest
- Submit for Delivery

What is it?



- Evolving technology still considered Early days
- Strongly positioned
- Extremely platform agnostic
- Targeted Deployments
- Worth the learning curve

What is it?




Kubernetes is Evolving
FAST!

Kubernetes

Micro Service Apps
Orchestrator



Team

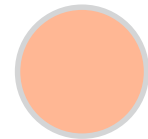


S1

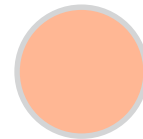
S2

S3

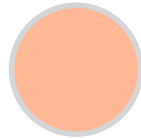




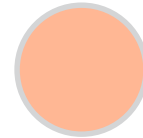
HTTPS



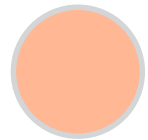
HTTPS



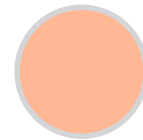
Search



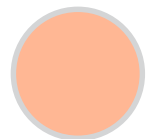
Auth



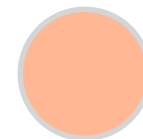
Key Vault



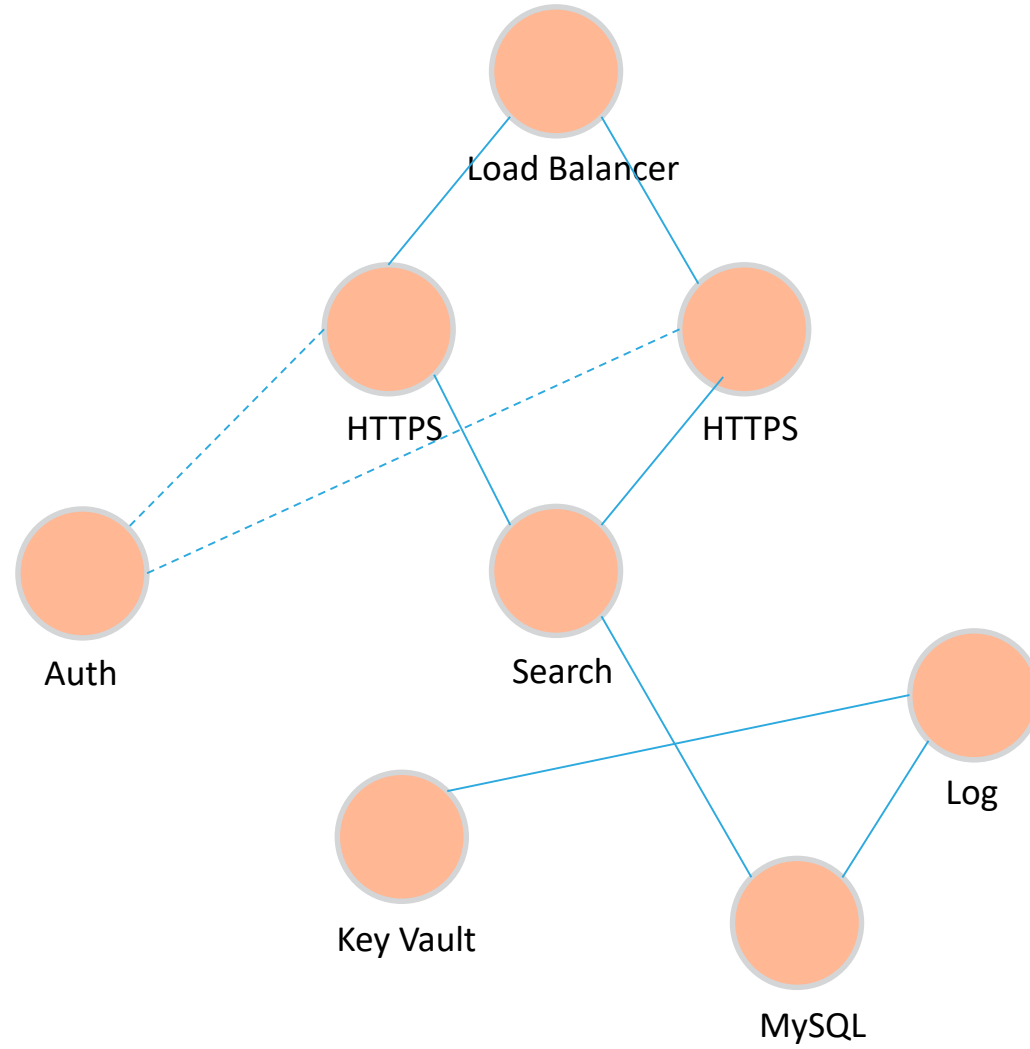
MySQL

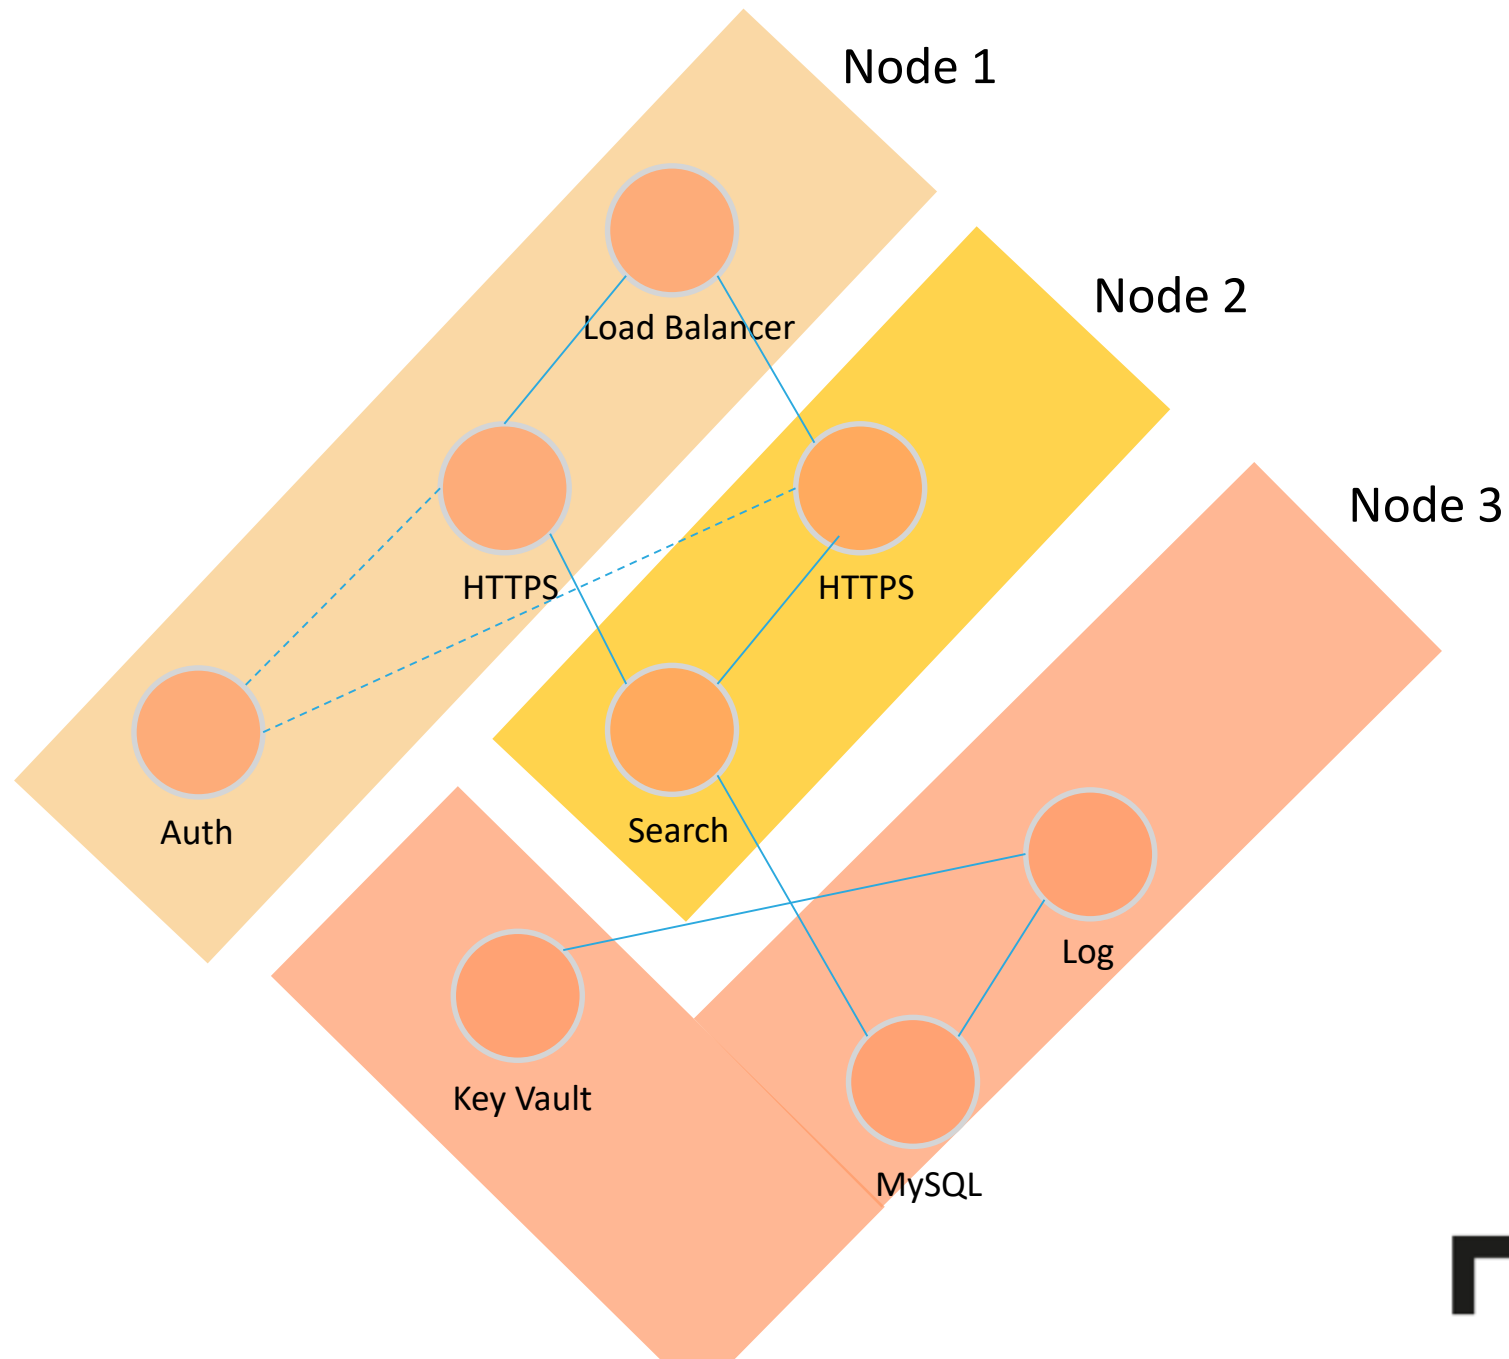


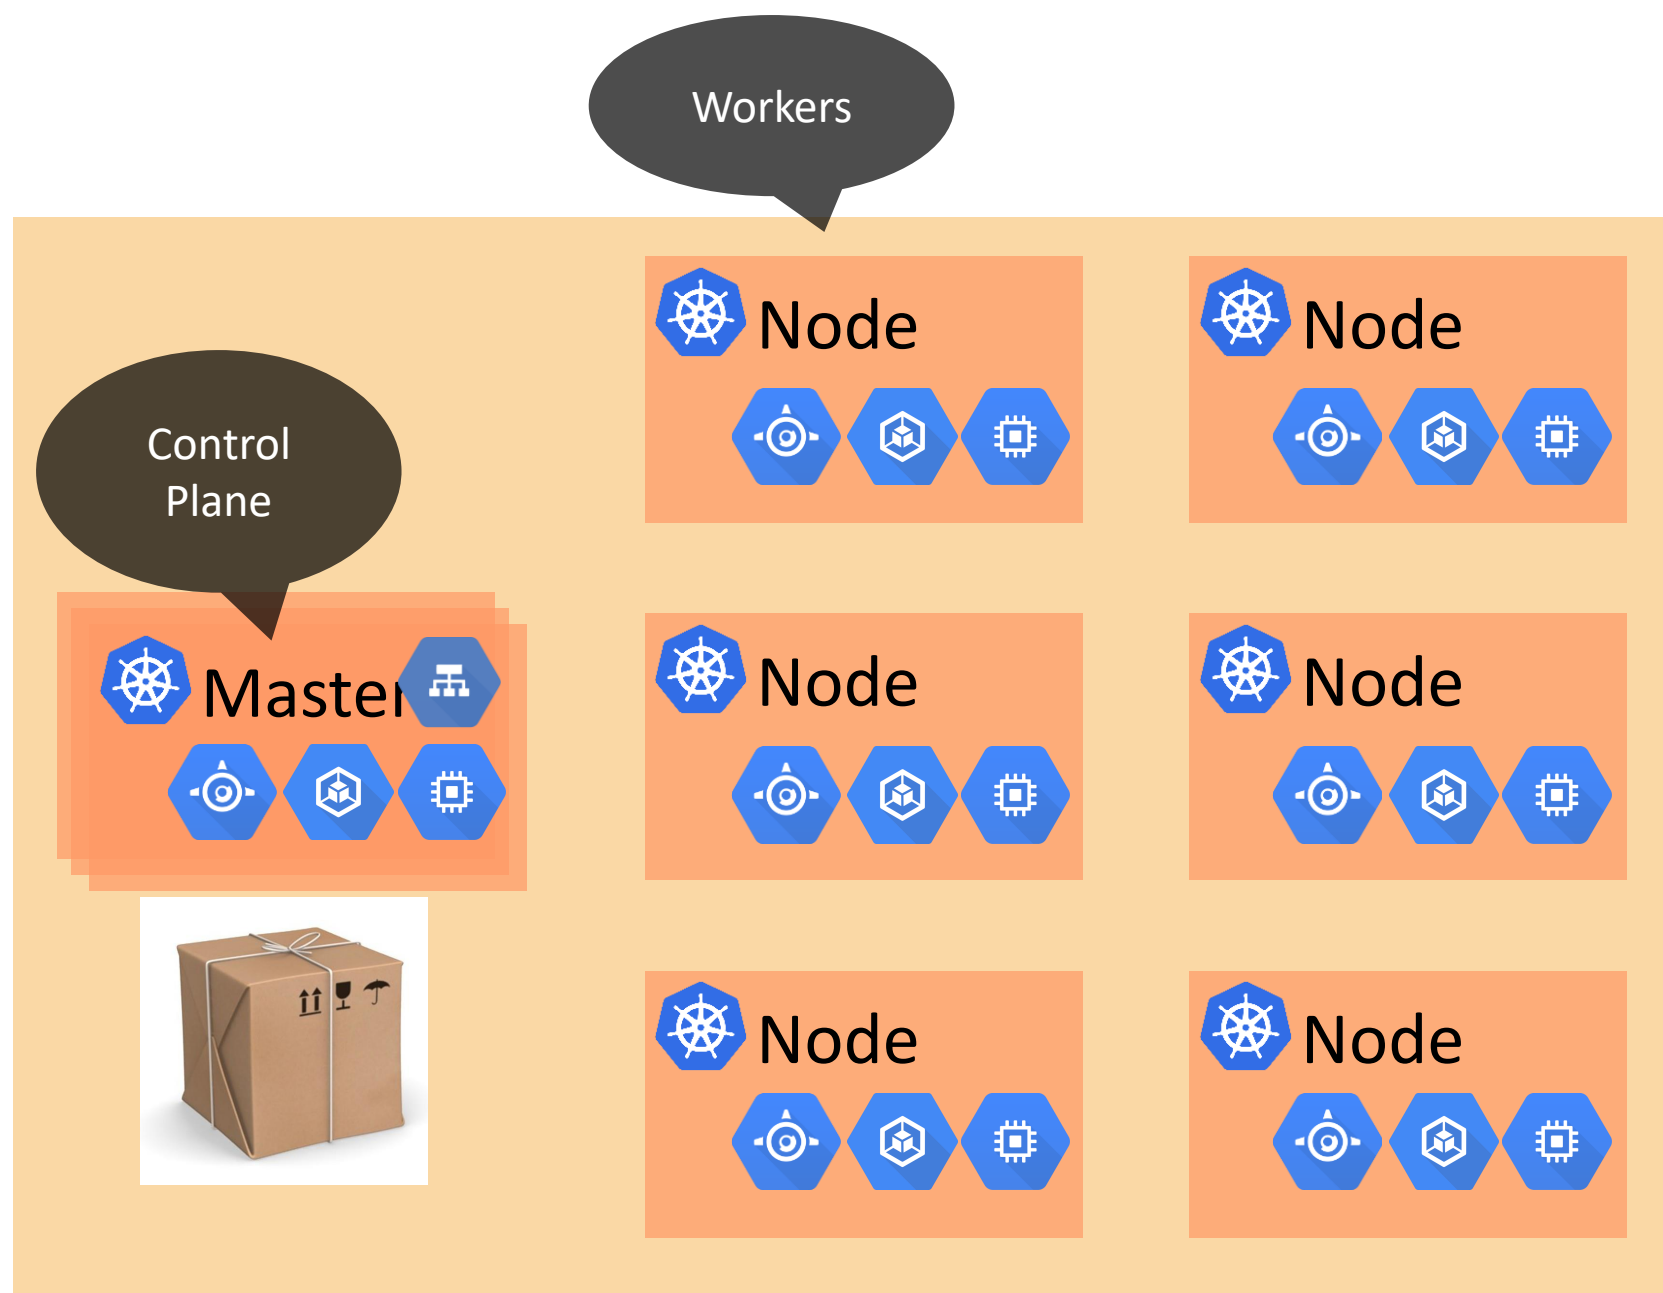
Log



Load Balancer

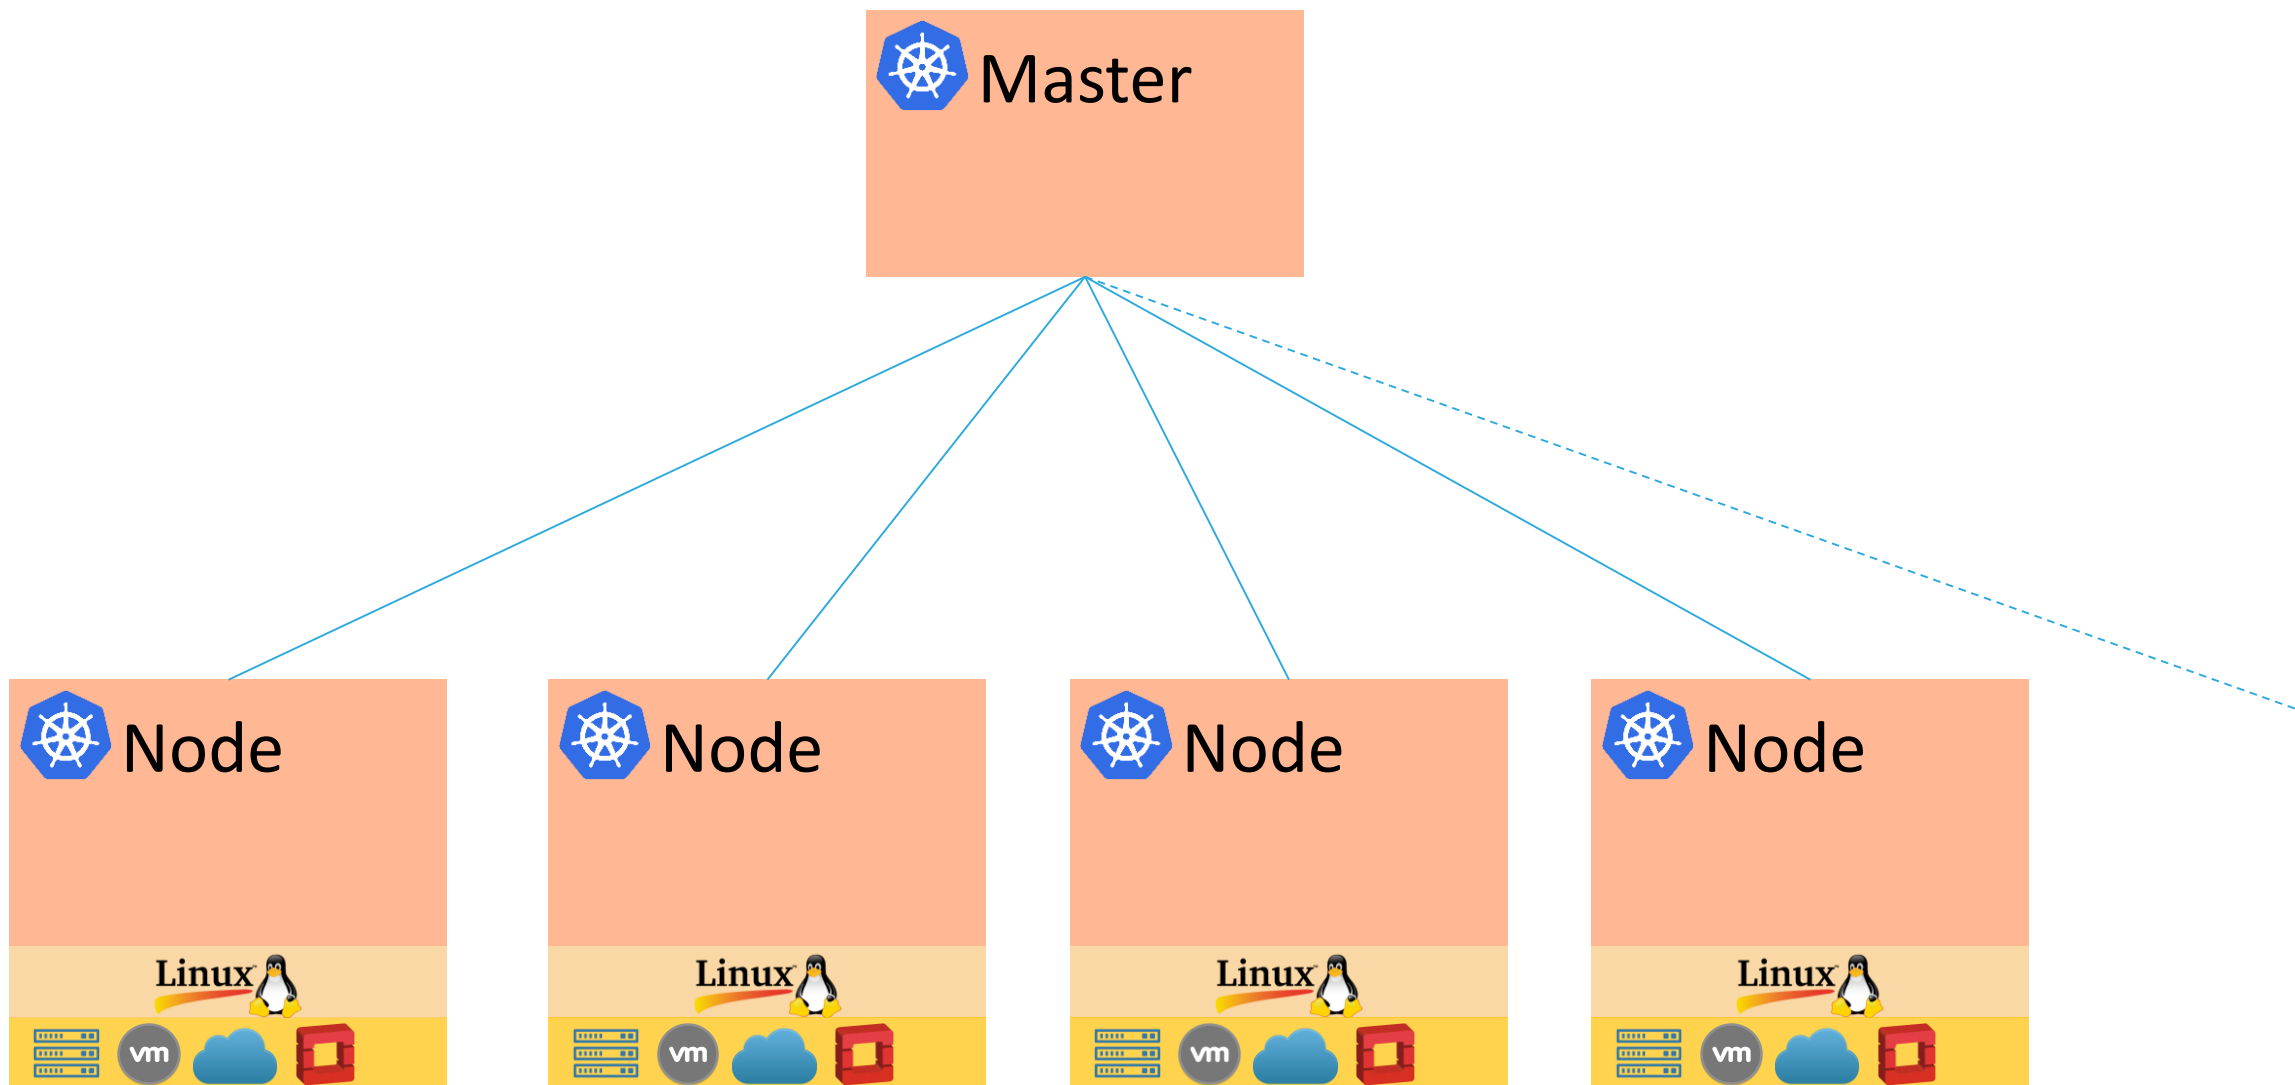


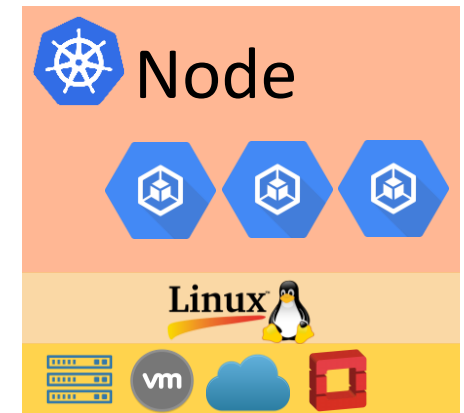
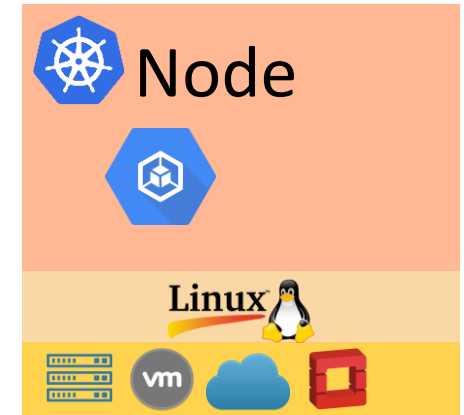
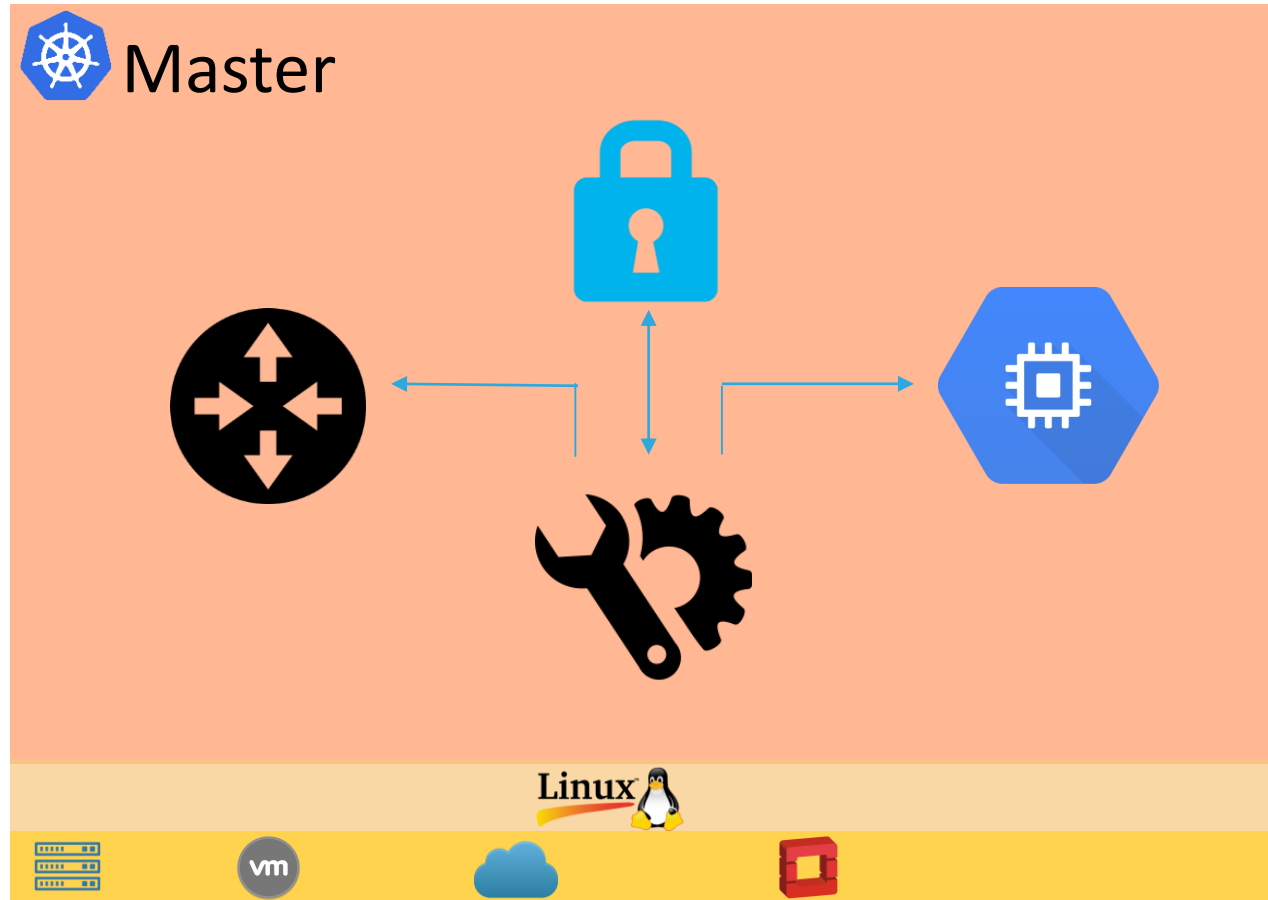
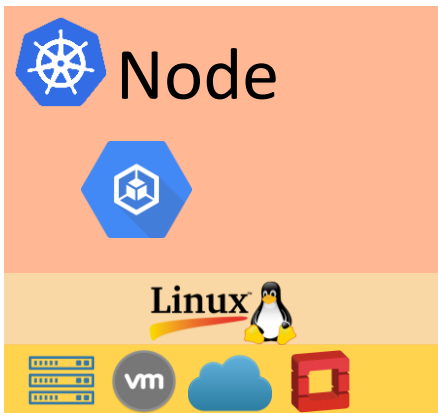
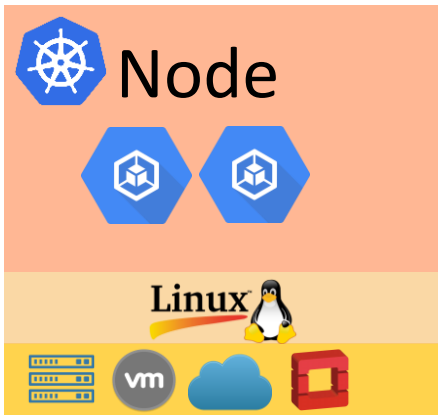


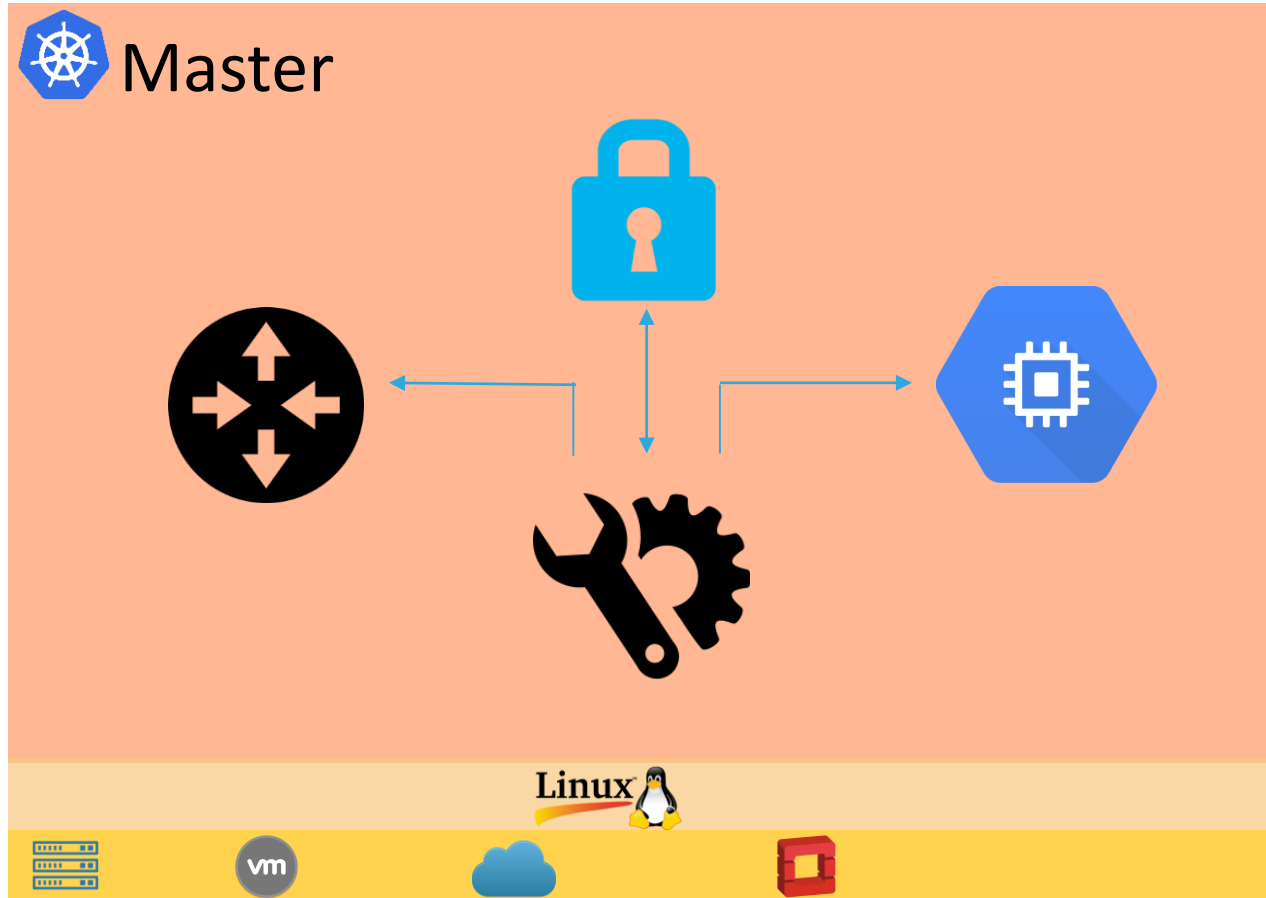


Masters

The Kubernetes Control Plane



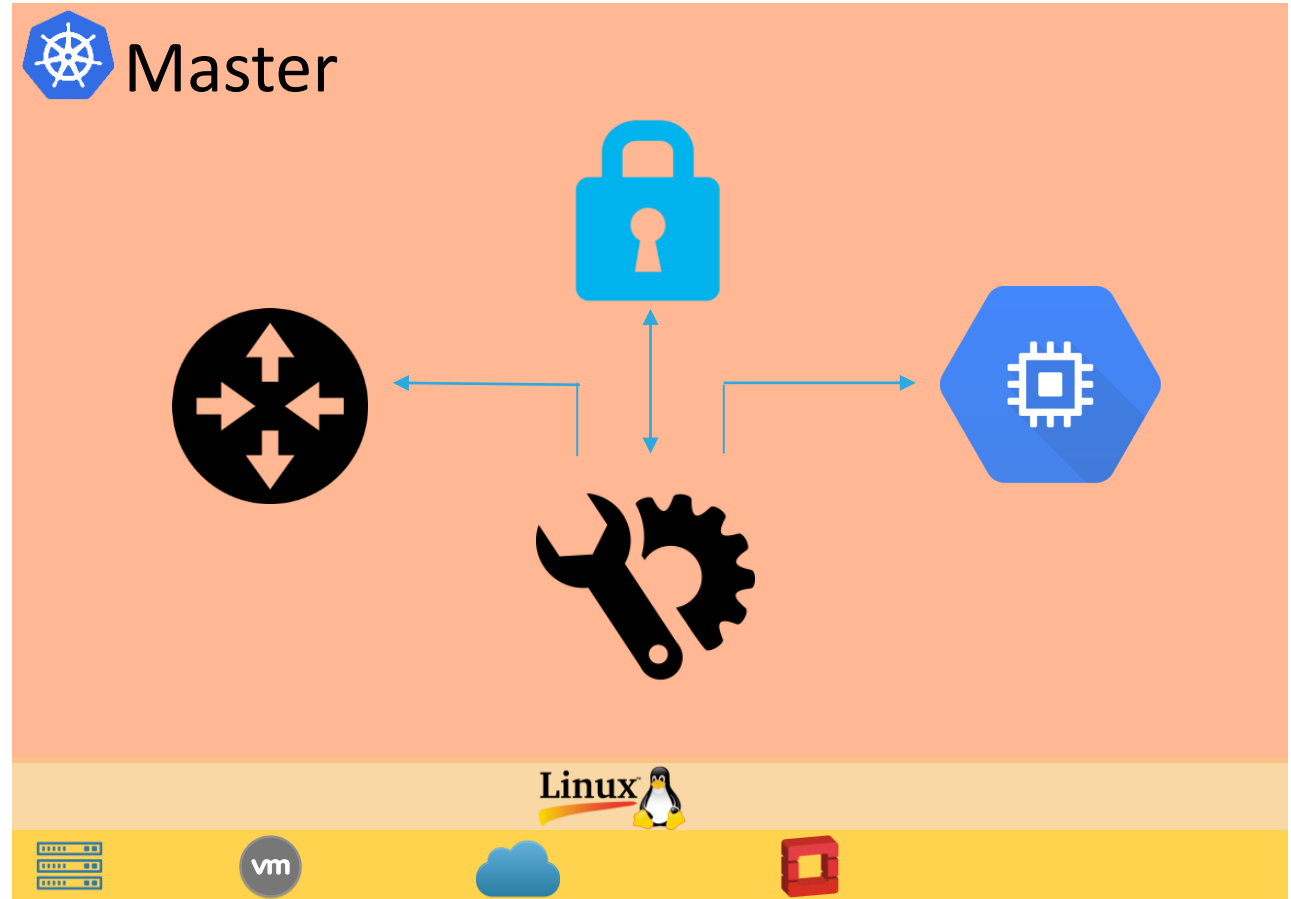




Kube-apiserver

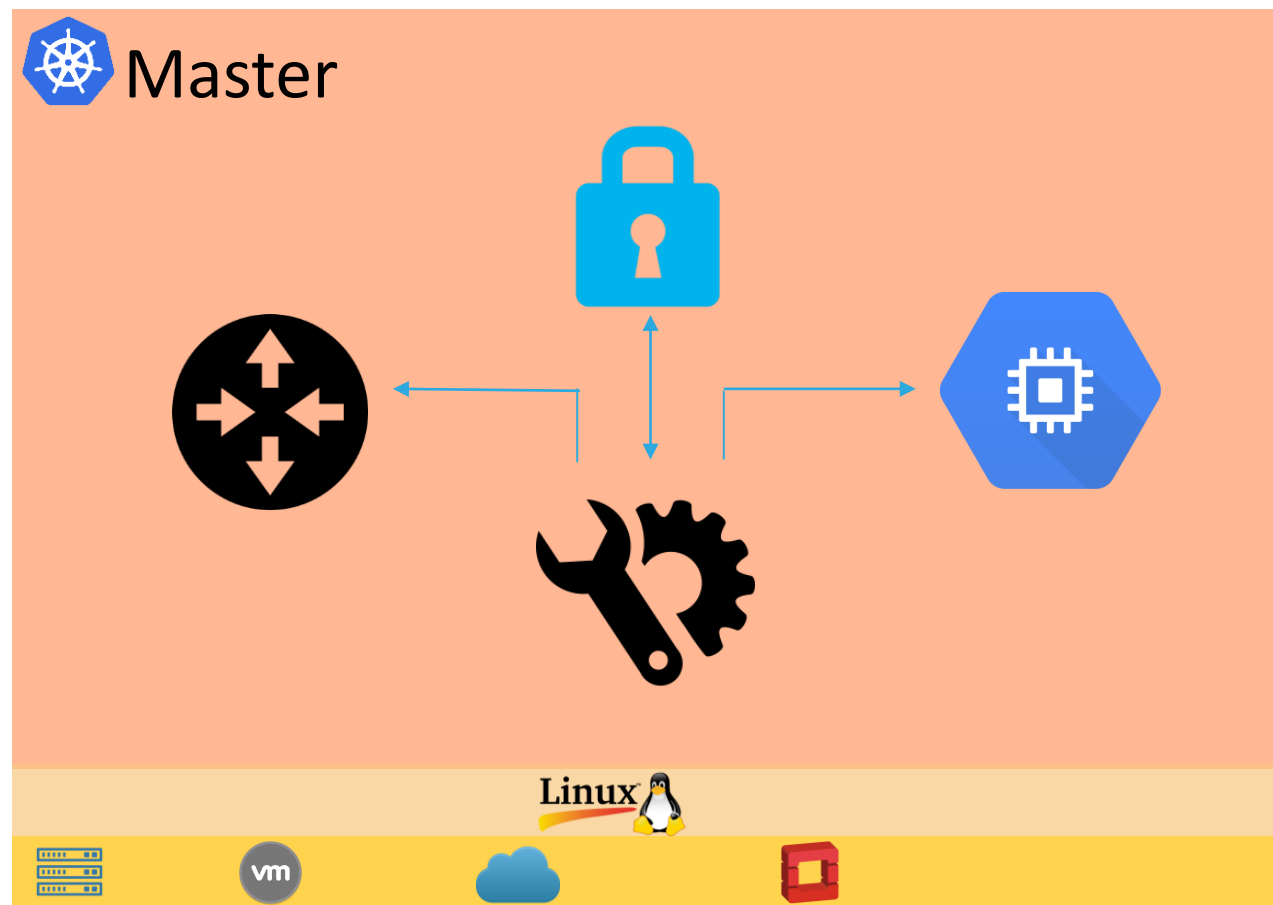


- Control Plane Front End
- Exposes ReST API
- Processes JSON Manifest Files



Cluster Store

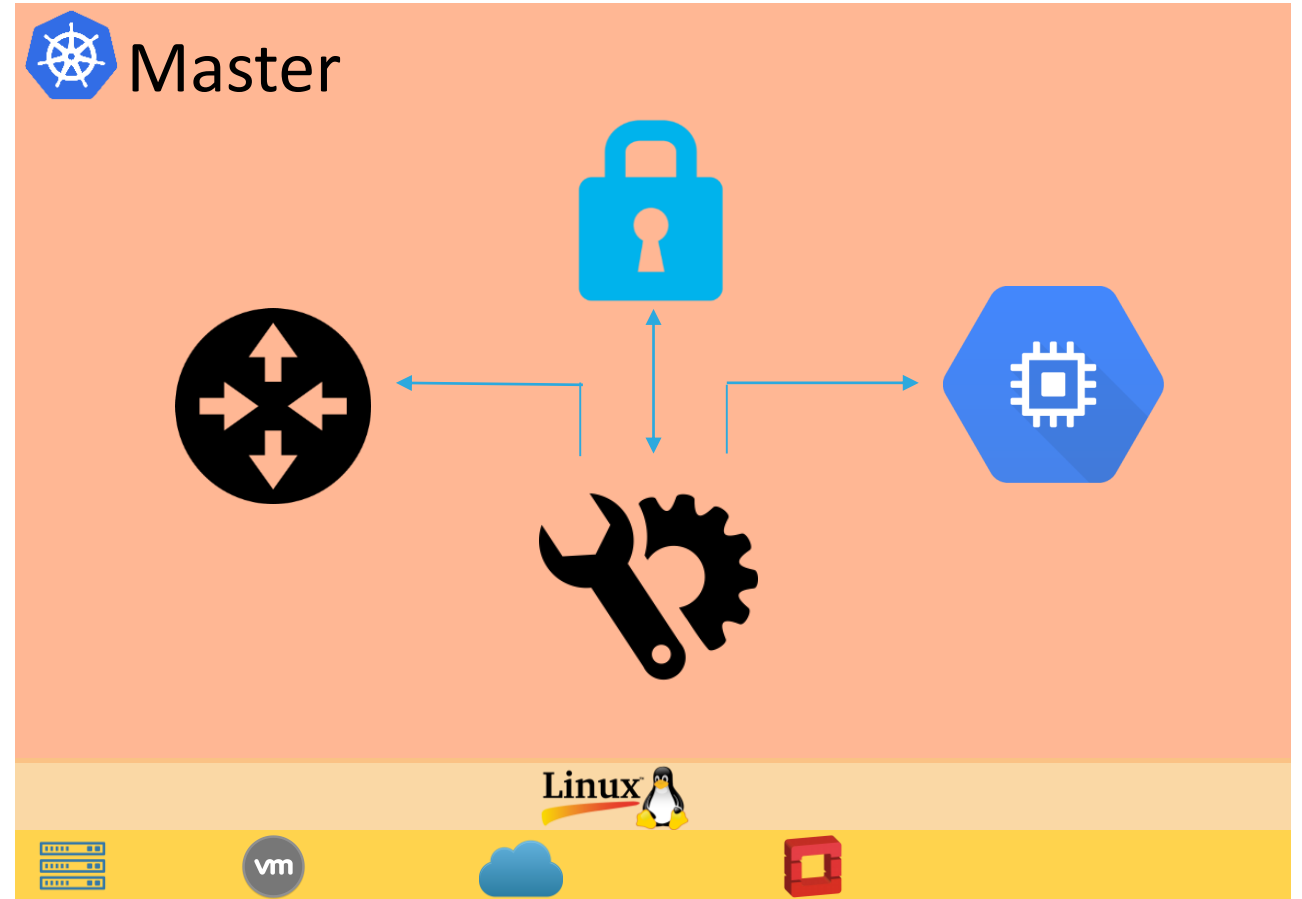
- Persistent Storage
- Cluster state and config
- Uses **etcd**
 - No SQL Database
 - Key Value Store
- Distributed, Consistent, Watchable
- Source of 'Truth' – Back It Up!



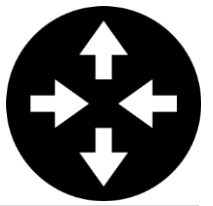
Kube-controller-manager



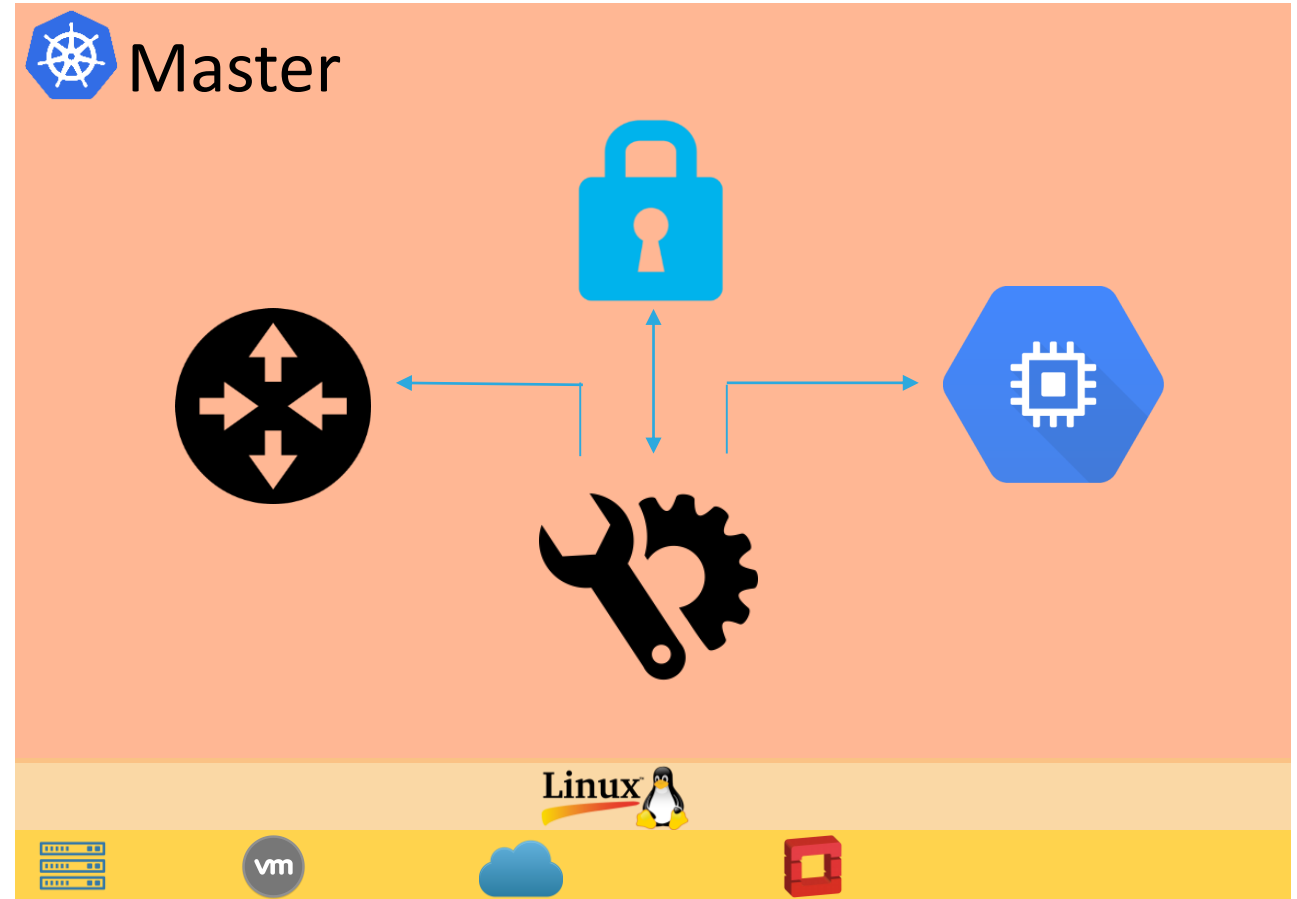
- Controller of controllers
 - Node Controller
 - Endpoint Controller
 - Namespace Controller
- Loops watching for changes
- Maintains desired state



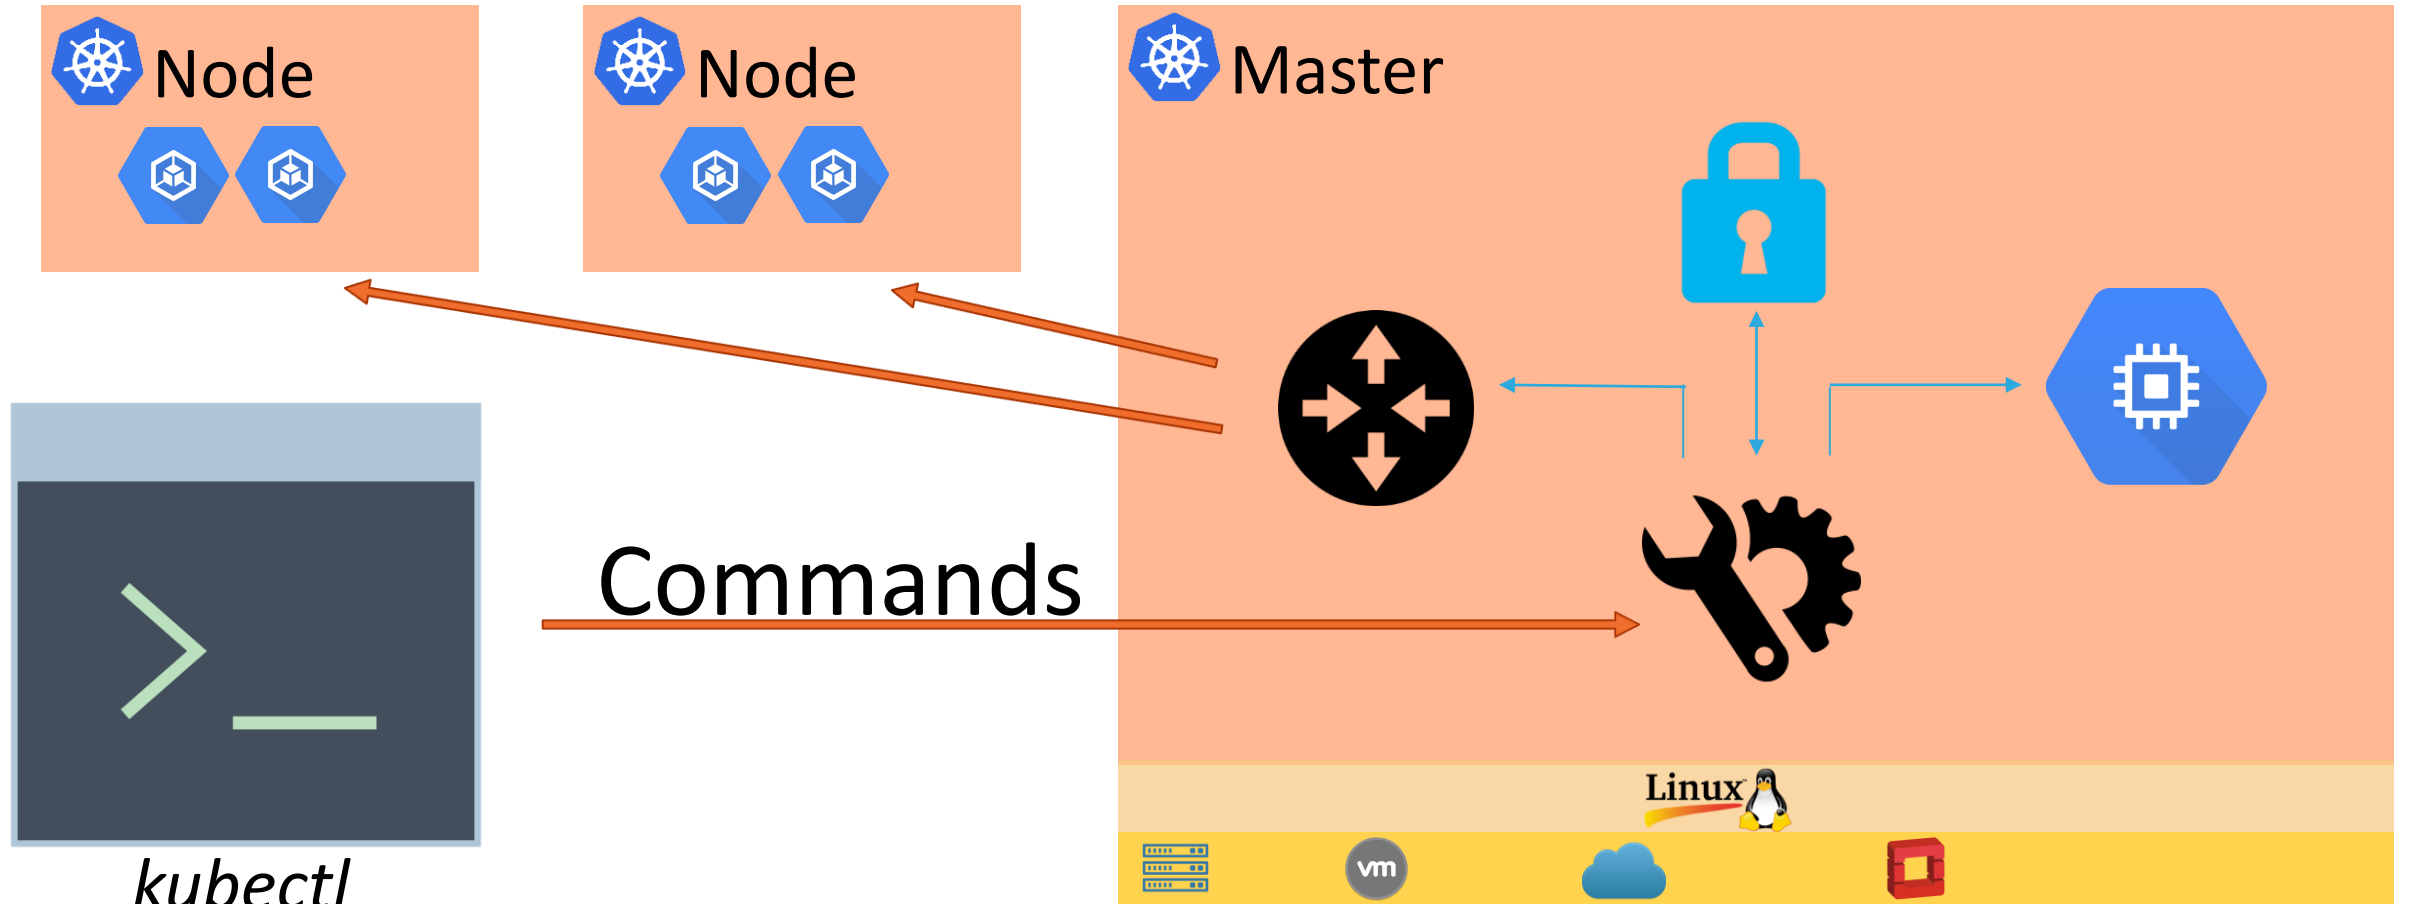
Kube-scheduler



- Watches **apiserver** for new Pods
- Assigns work to Nodes
 - affinity/anti-affinity
 - constraints
 - resources



Master



Nodes



The Kubernetes Workers

 Node



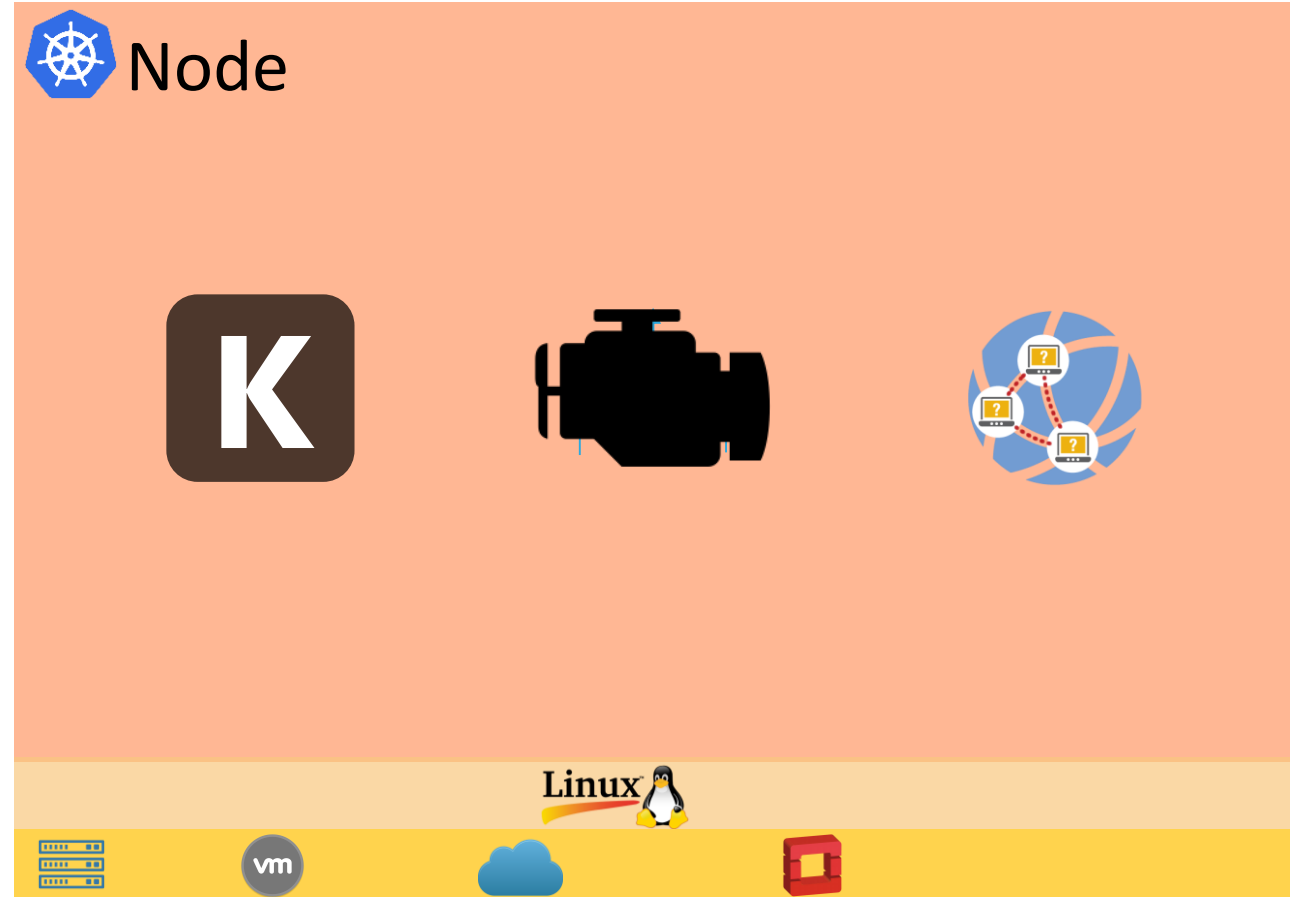
Linux 



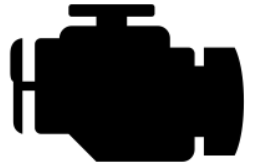
Kubelet



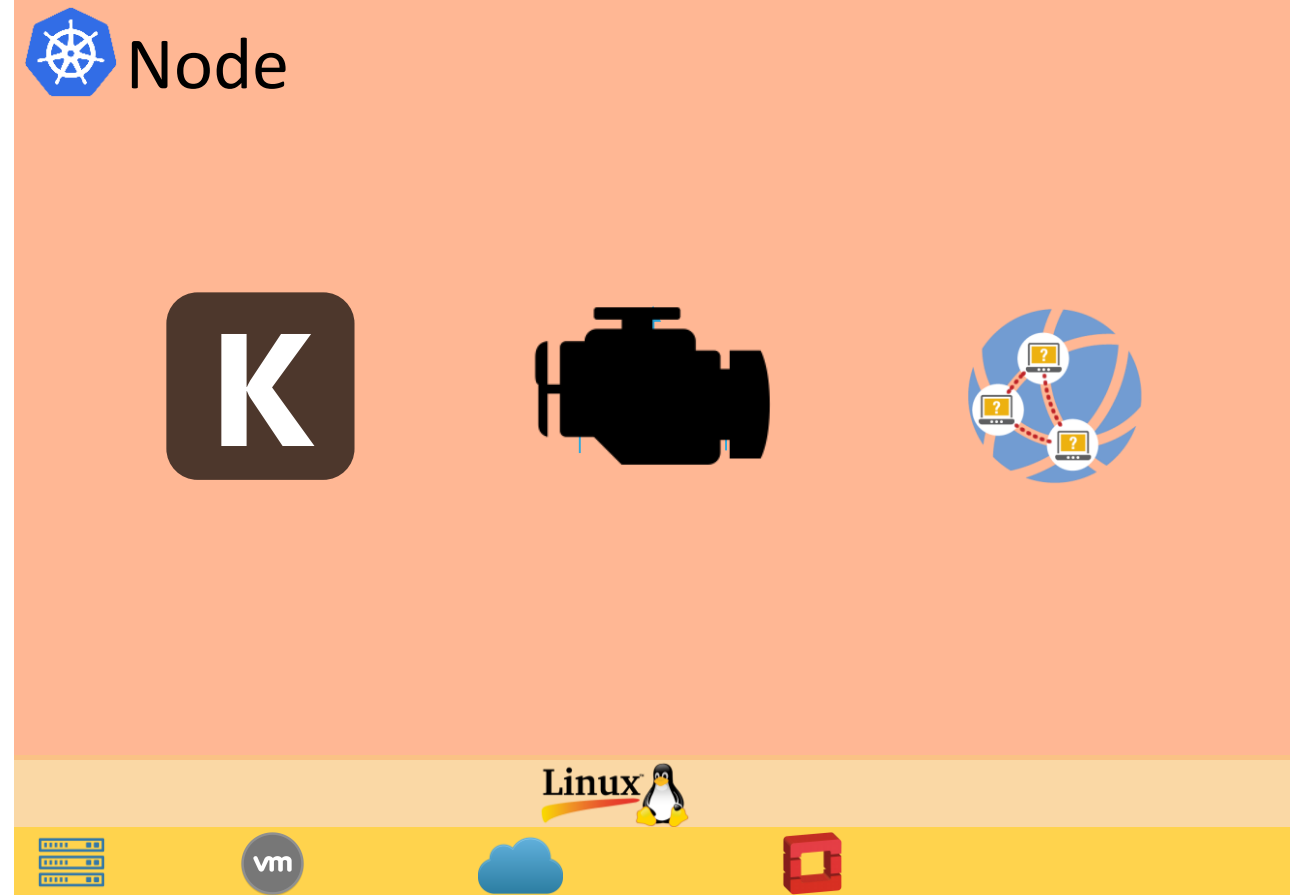
- Main Kubernetes Agent
- Registers node with cluster
- Watches **apiserver**
- Instantiates **Pods**
- Reports back to master
- Exposes endpoint TCP 10255
 - /spec
 - /healthz
 - /pods



Container Engine



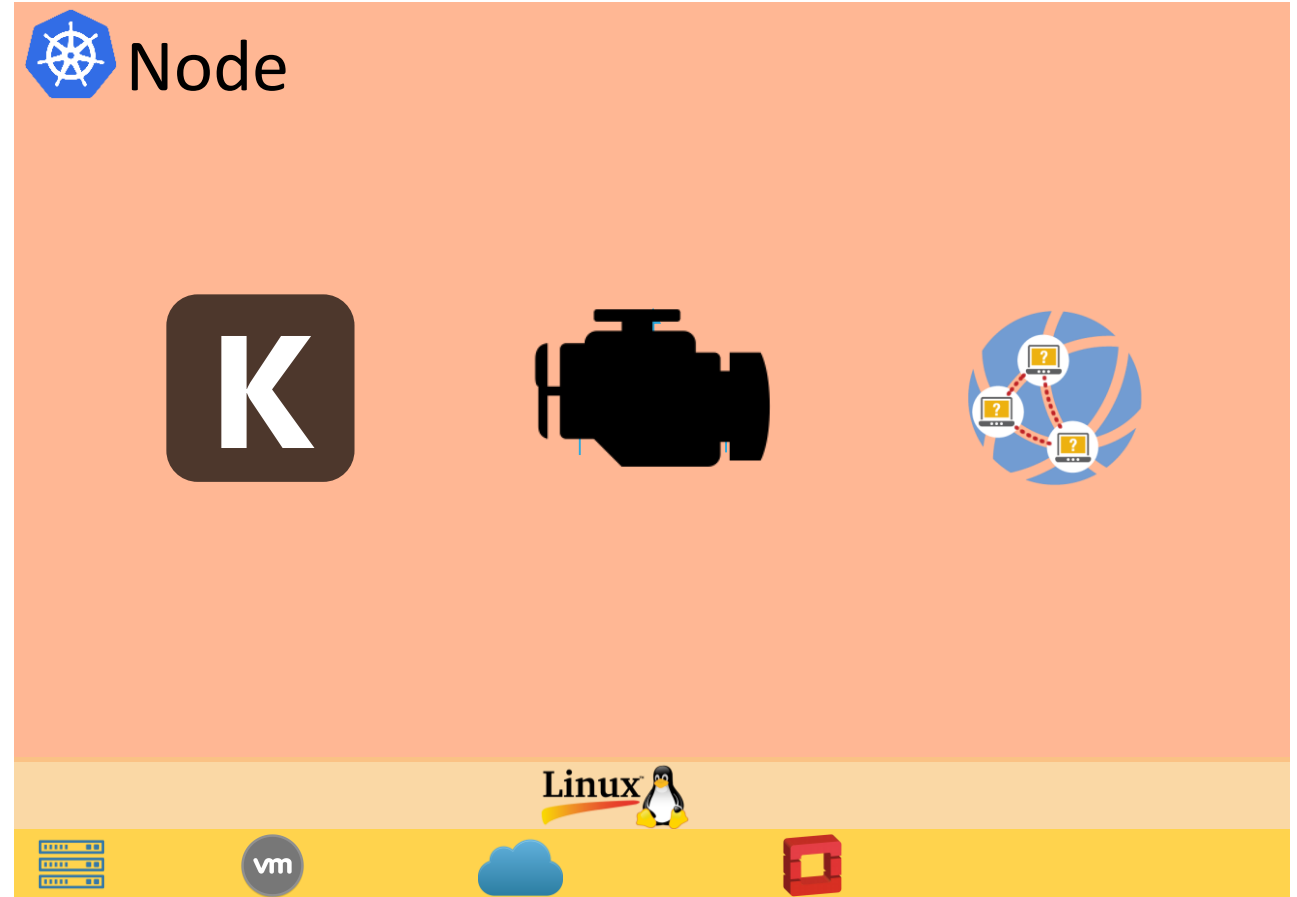
- Container Management
 - Retrieving Images
 - Starting & Stopping Containers
 - ...
- Pluggable
 - Docker
 - CoreOS rkt



Kube-proxy

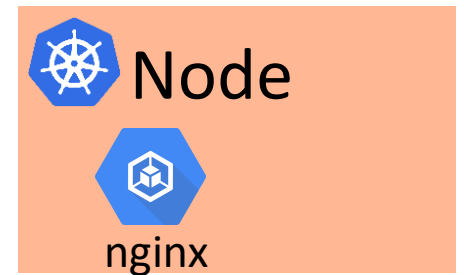
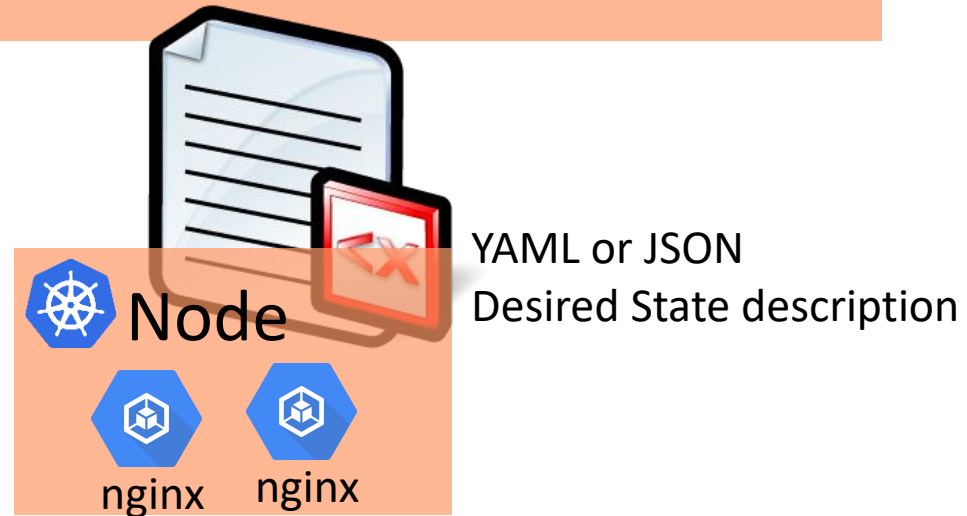
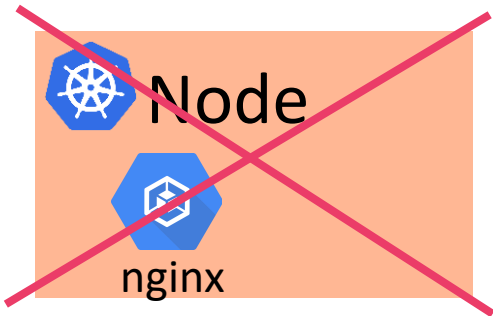
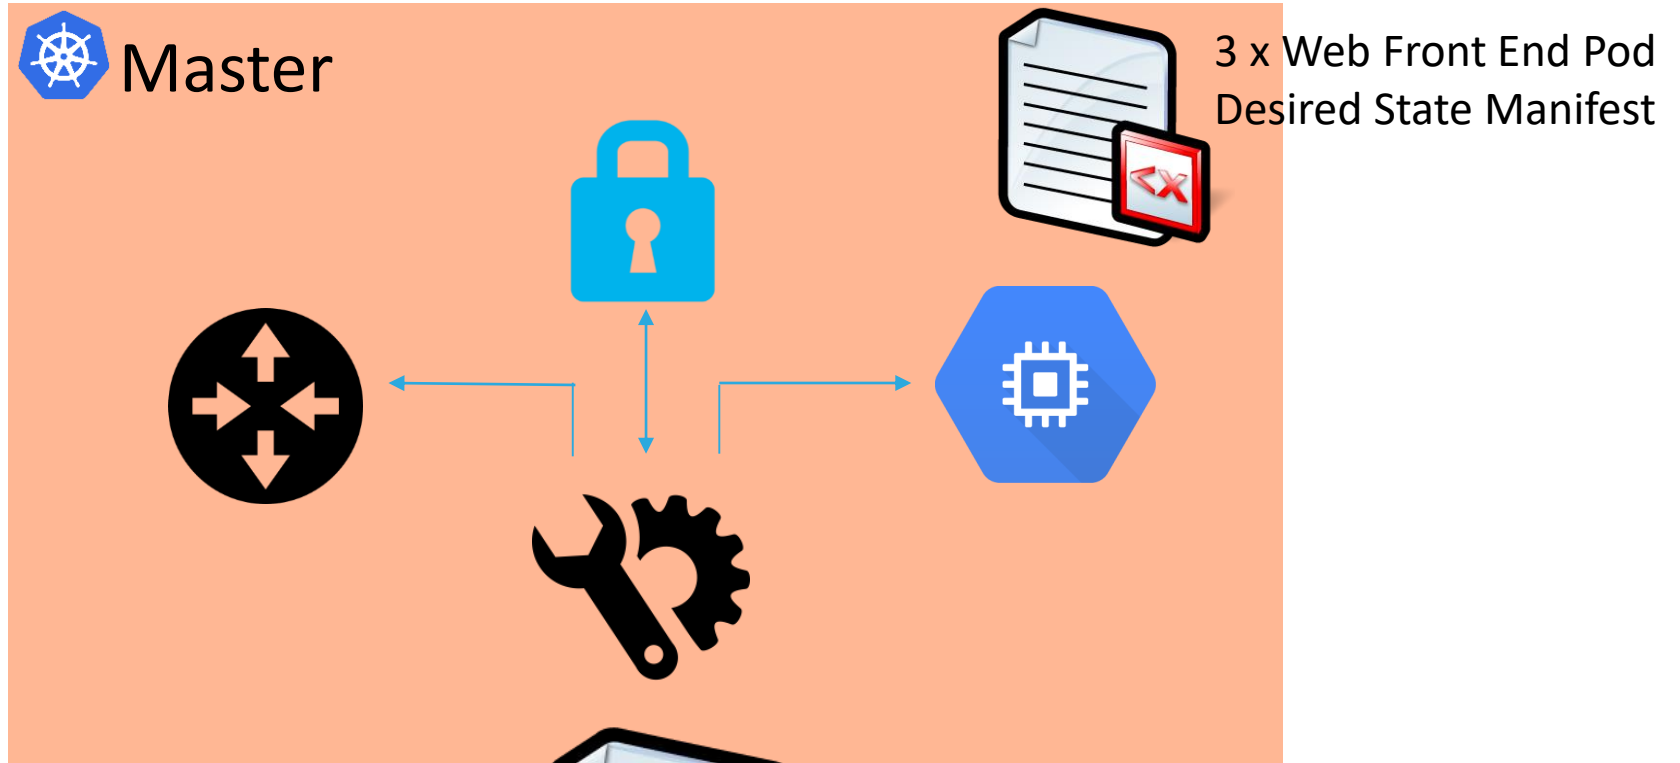


- Networking
 - Pod IP address
 - All containers in a pod share a single IP
 - Load balances across all pods in a service



Declarative Model

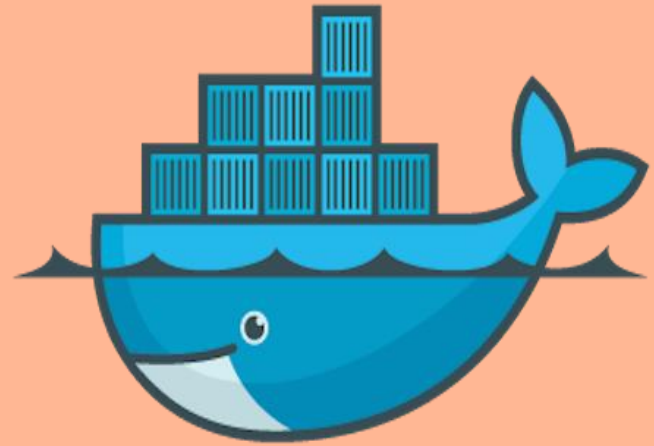
Desired State



Pods

Atomic Units

Atomic Unit Of Scheduling



docker



Containers



- Kubernetes runs container
- Only Inside Pods
- Pods may contain multiple containers

POD



POD



POD



POD



Pod



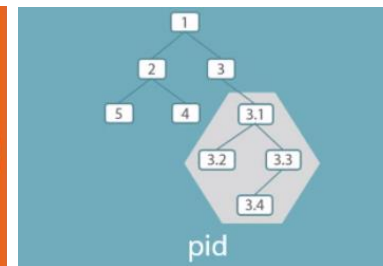
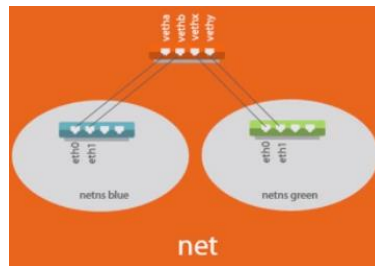
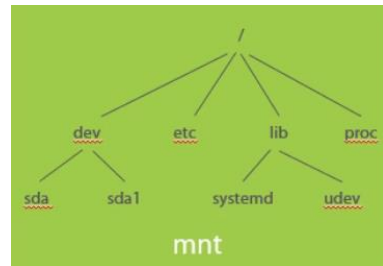
- Ring fenced environment
 - Network Stack
 - File System
 - Kernel Namespaces
 - ...
- Hosts Container(s)
- All Containers SHARE the Pod environment



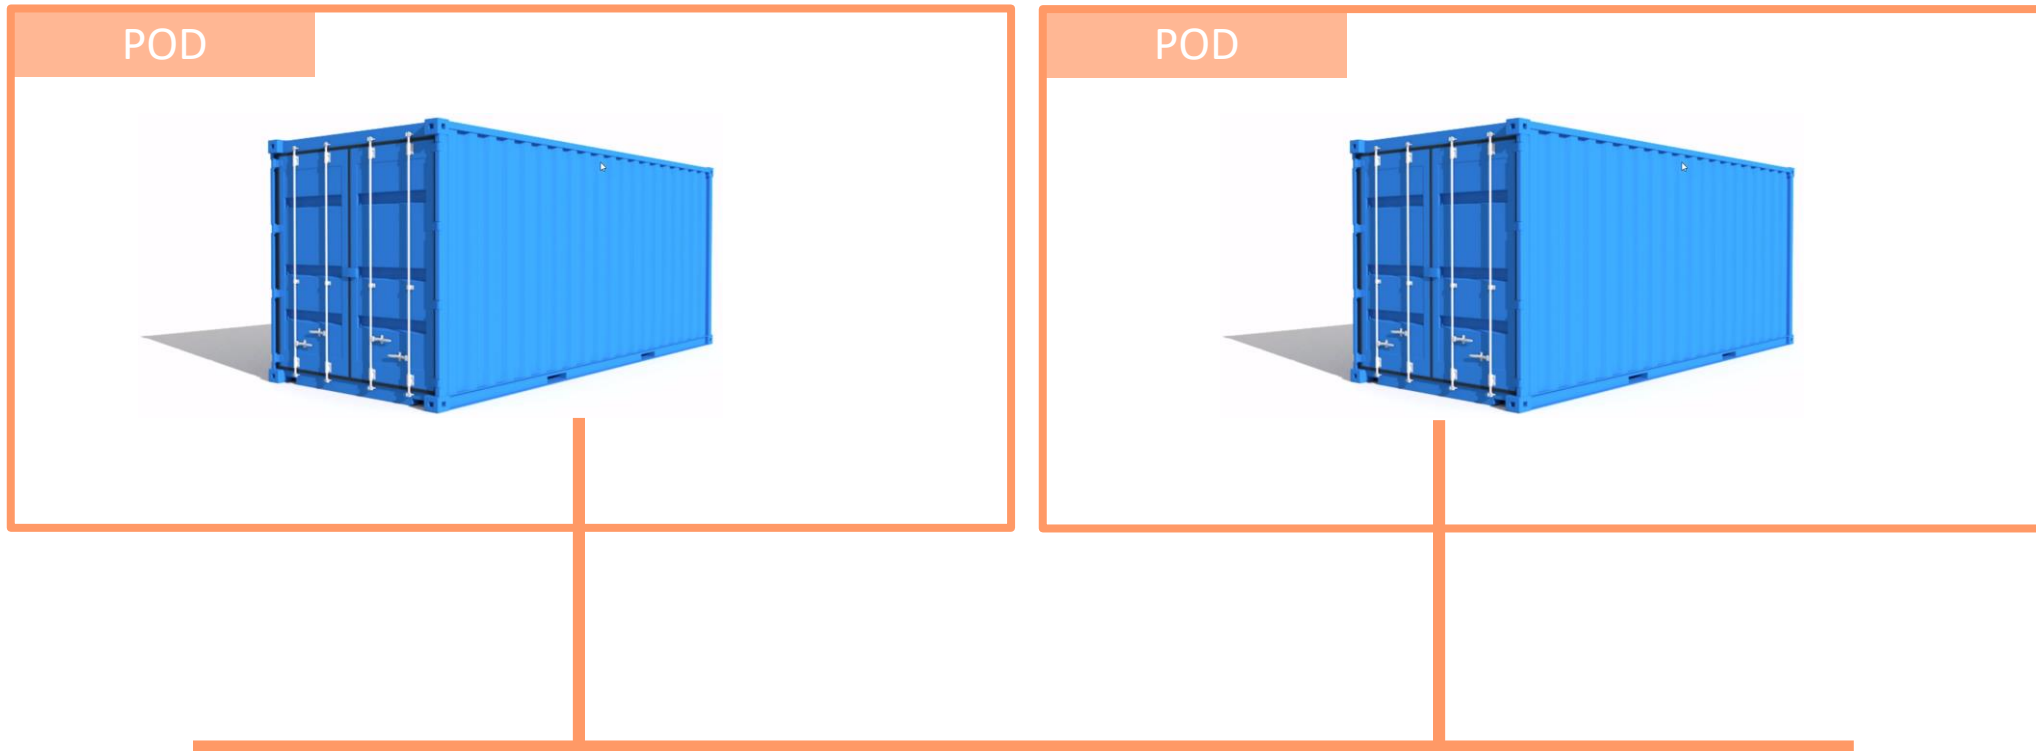
Tight Coupling



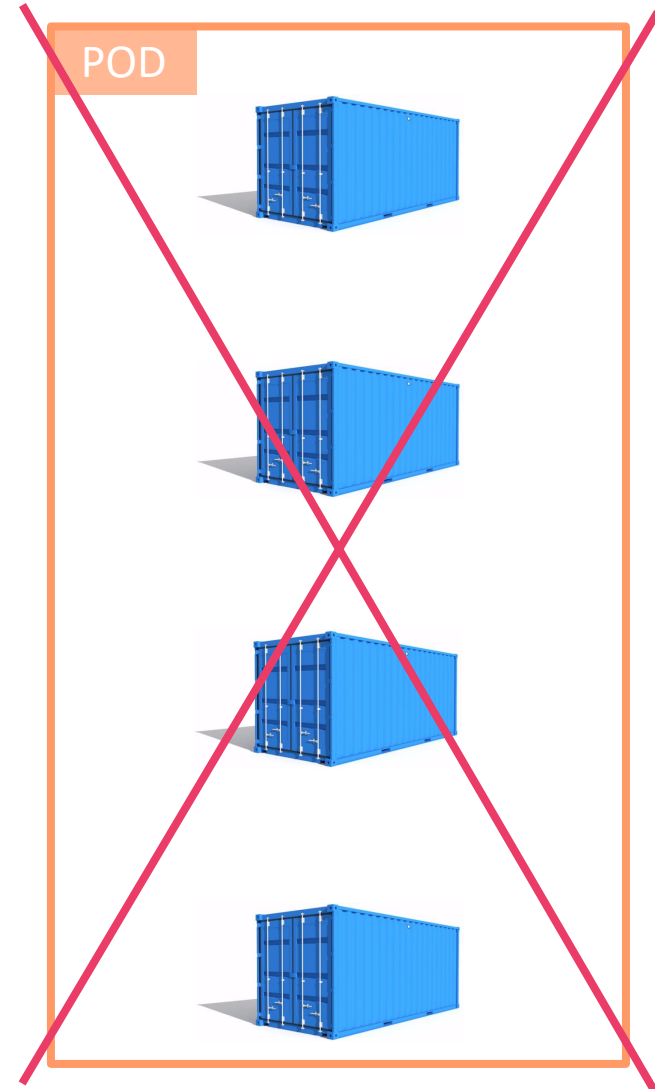
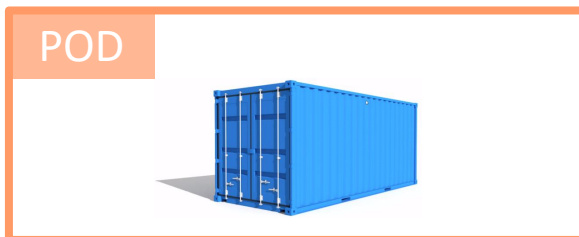
POD



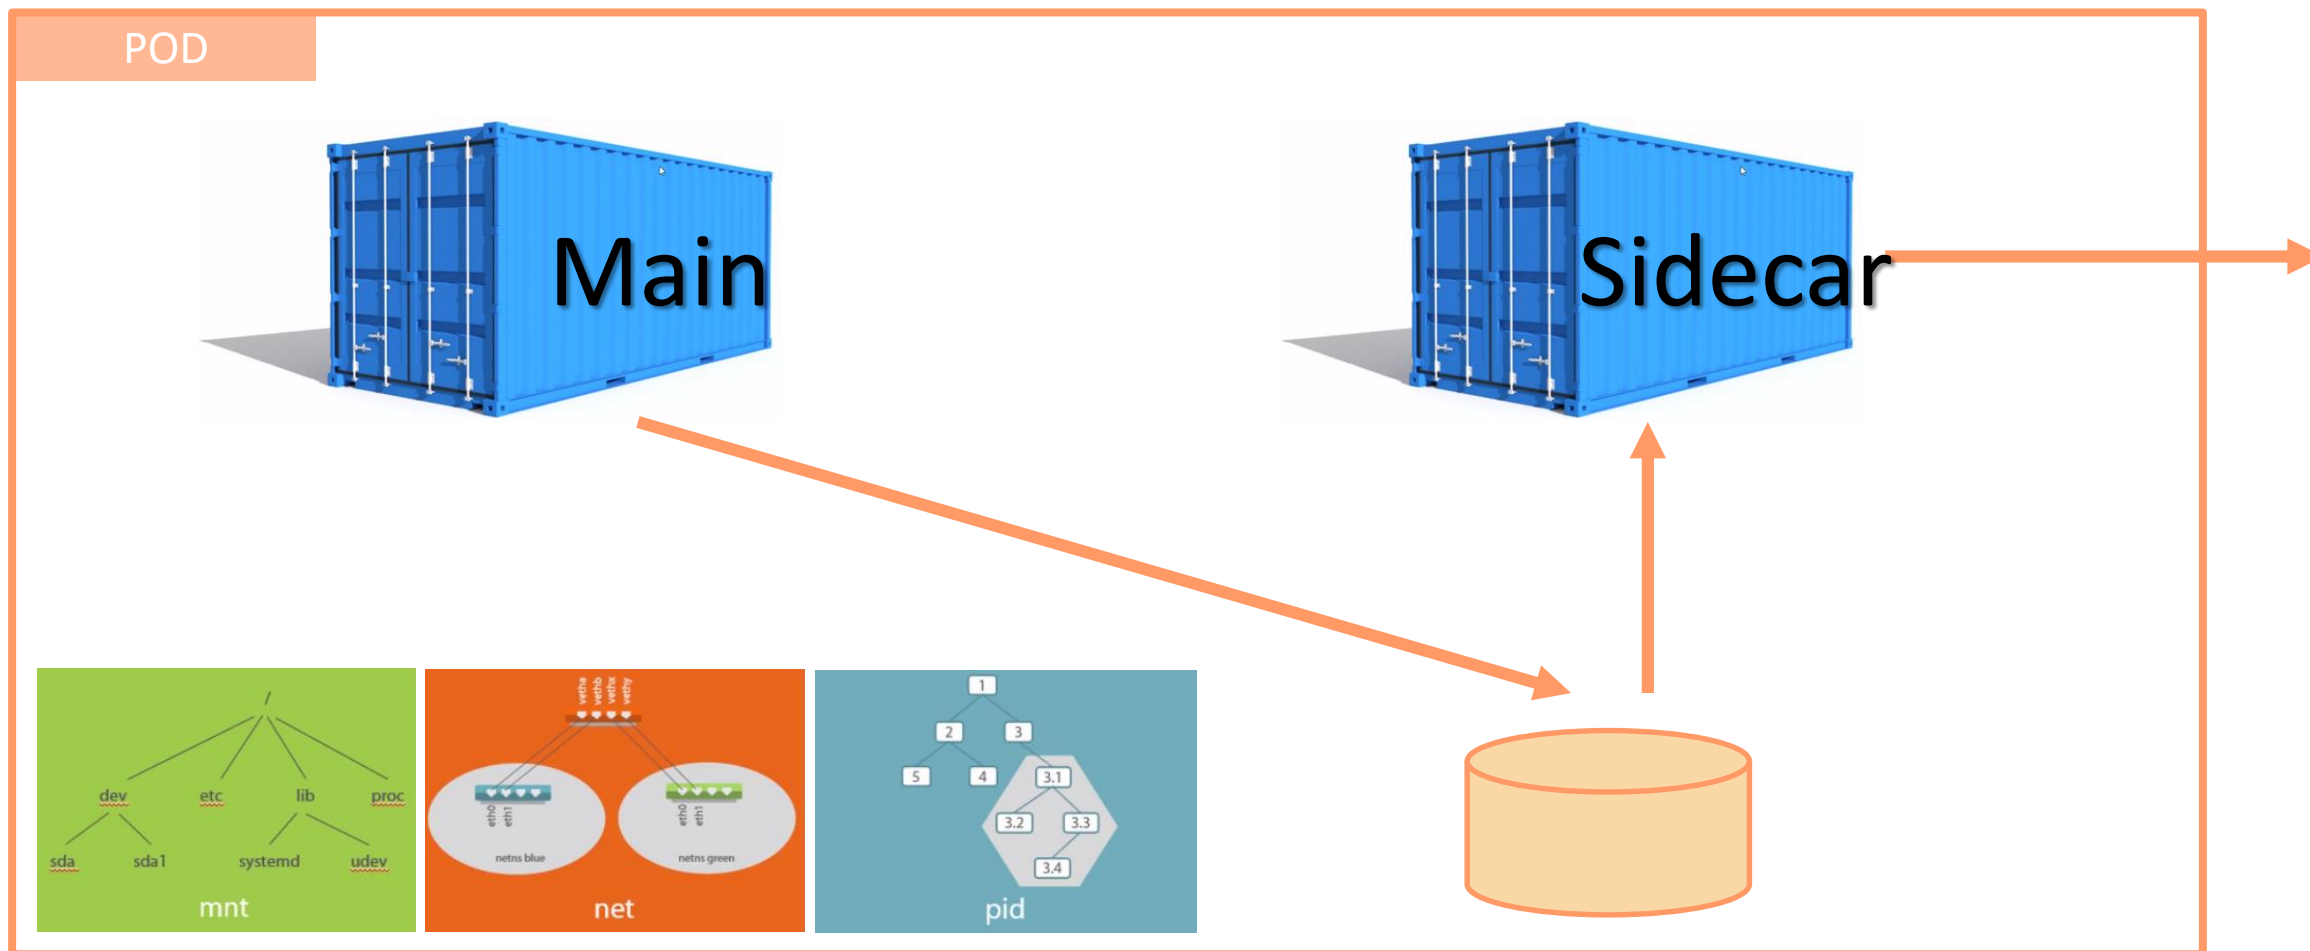
Loose Coupling



Scaling



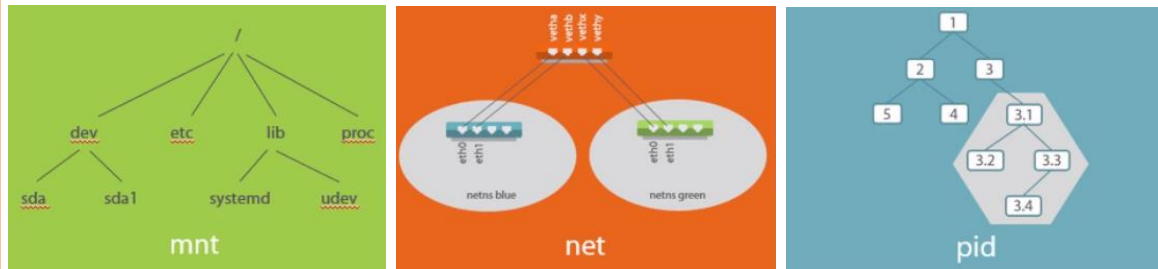
Multiple Containers



Atomic

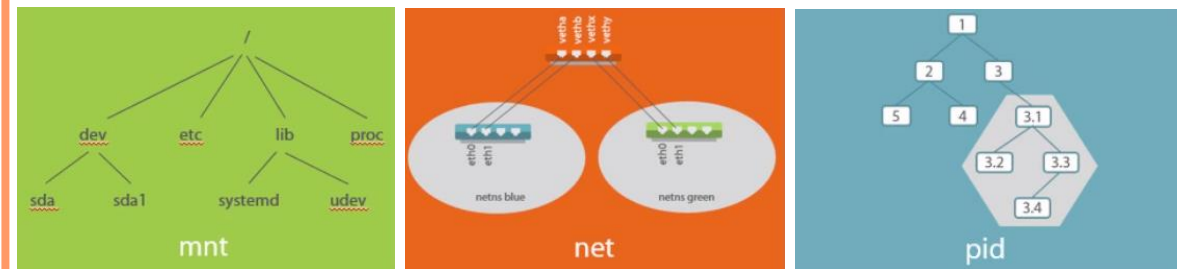


POD



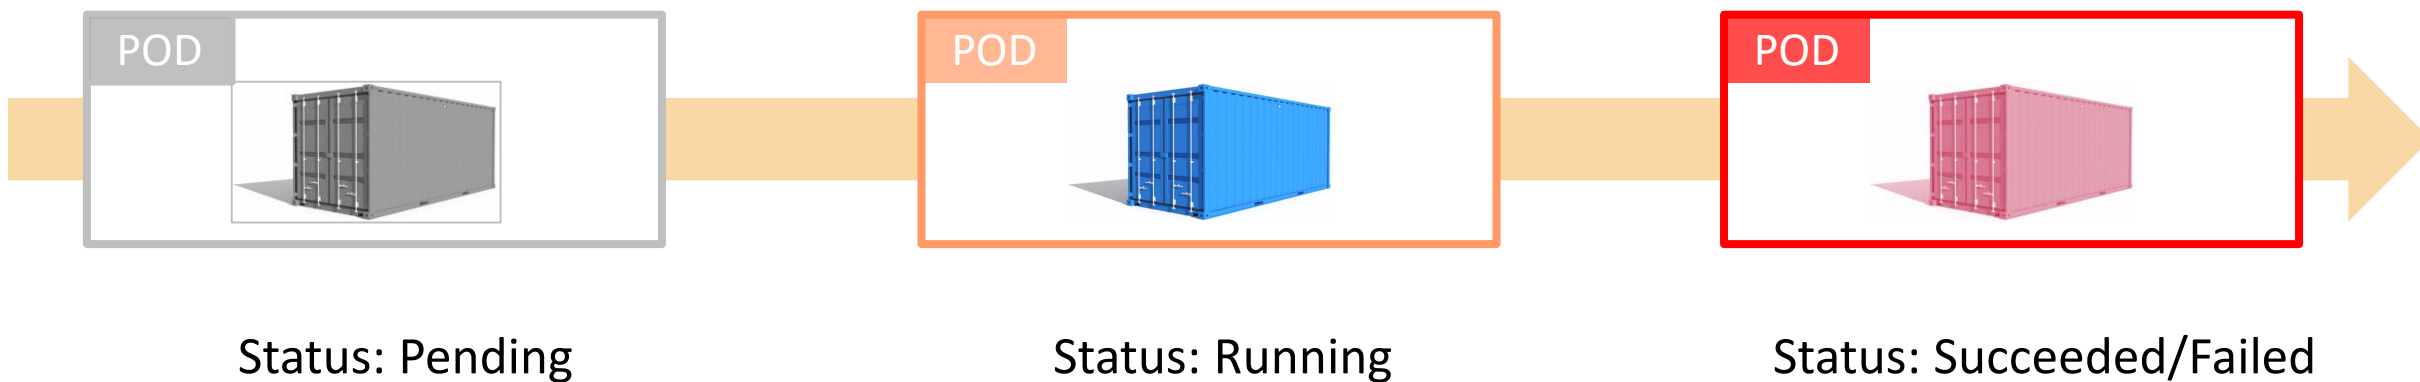
Status: Ready

POD

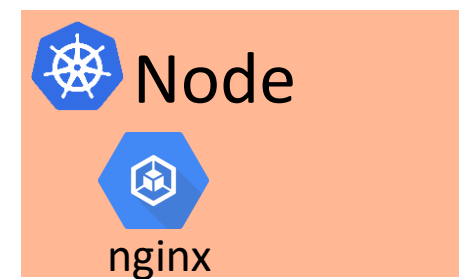
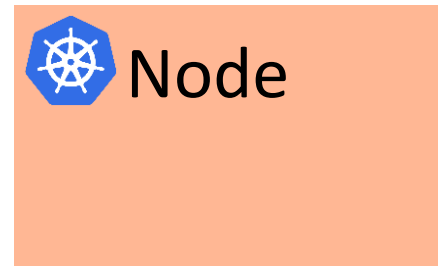
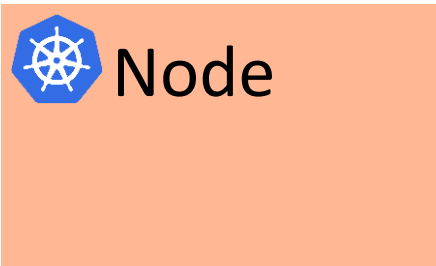
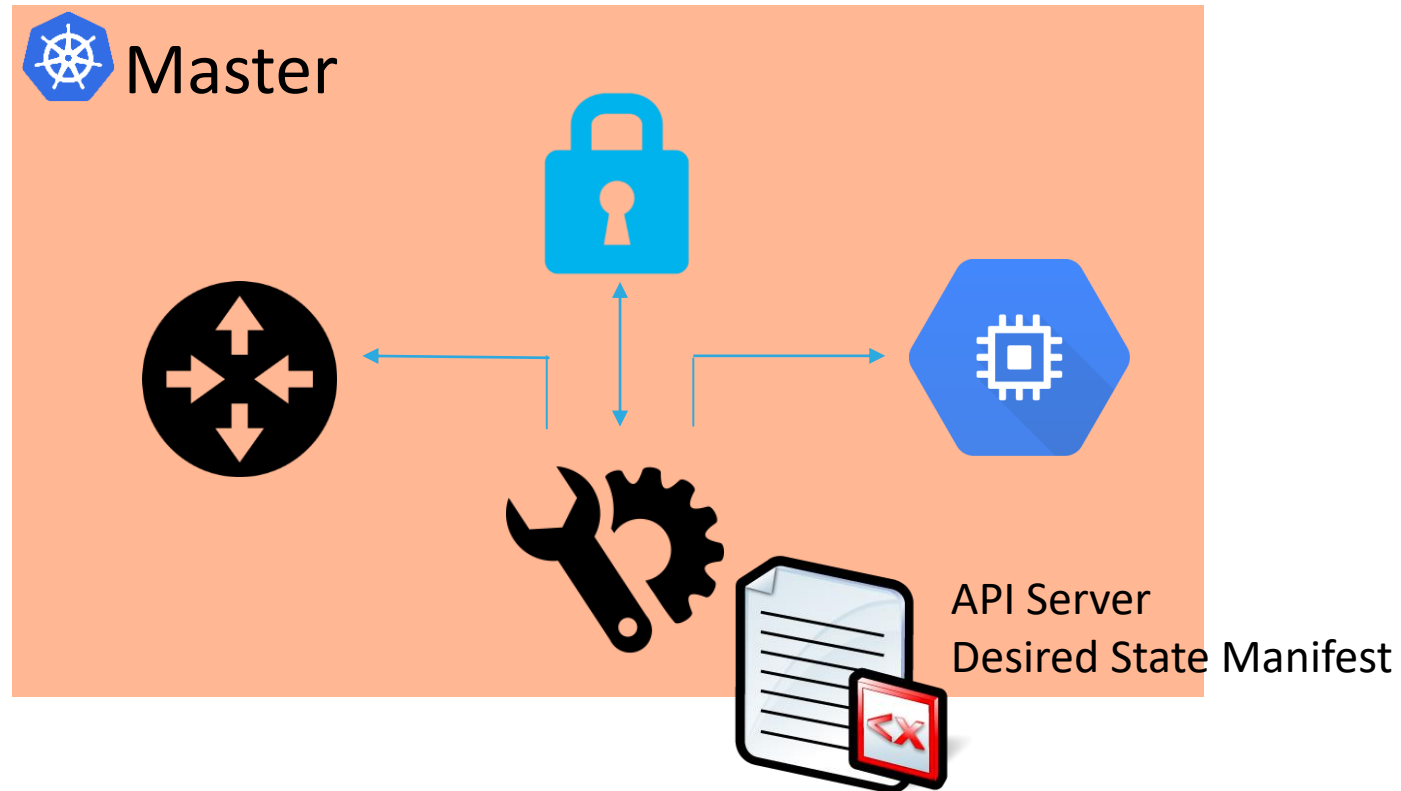


Status: Pending

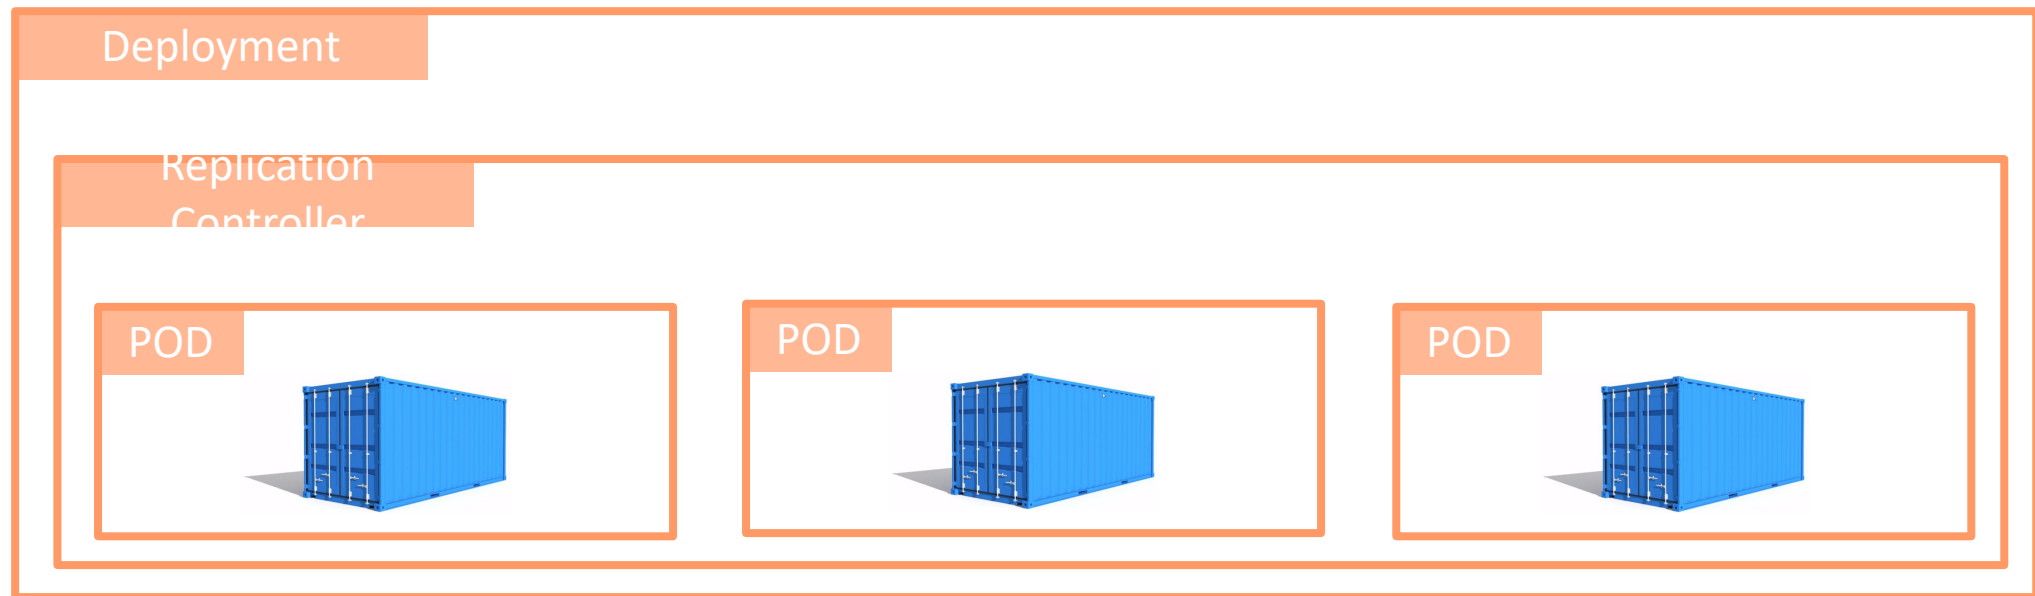
Pods are Mortal!



Deployment



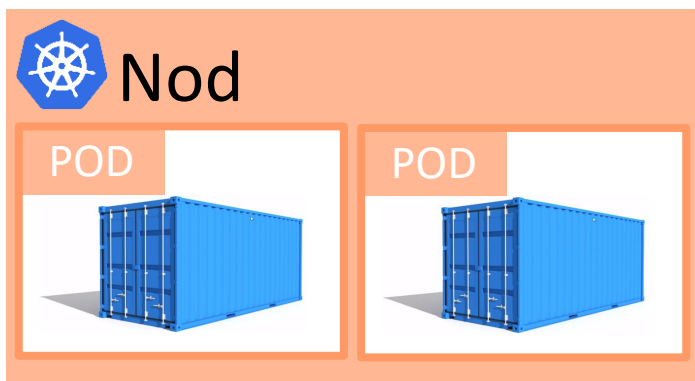
Deployment Abstraction



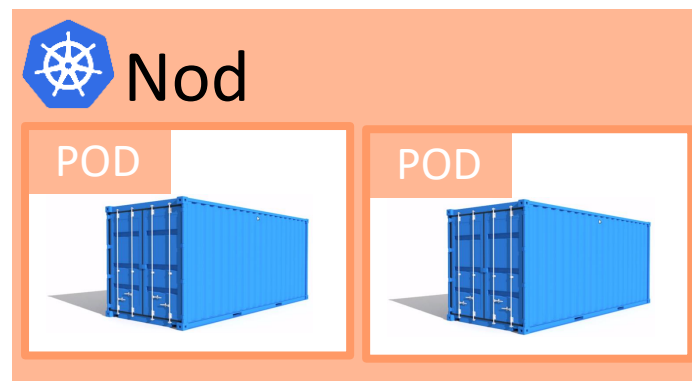
Services

Atomic Units

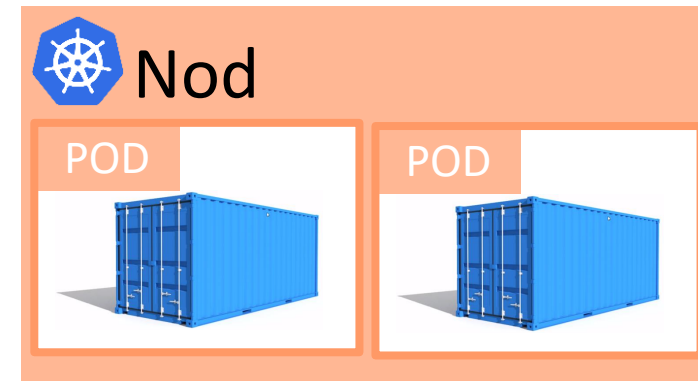
IP Churn



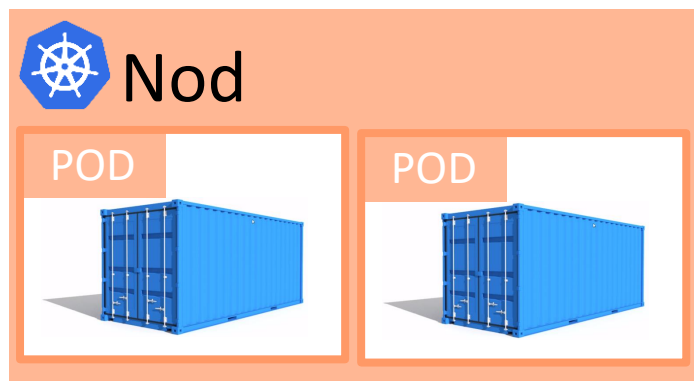
172.16.1.45 172.16.1.66



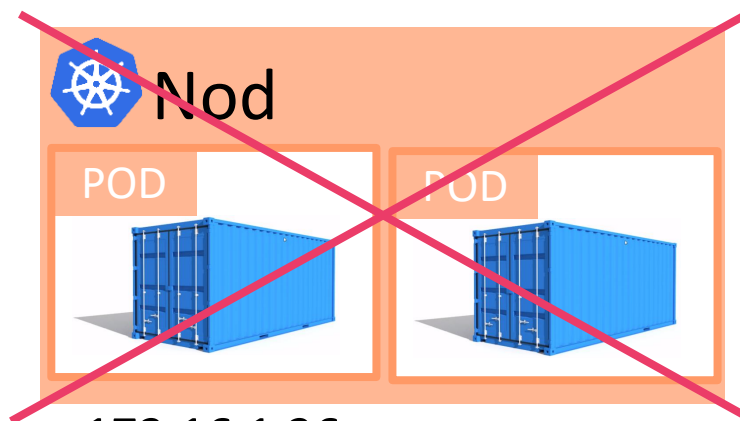
172.16.1.72 172.16.1.75



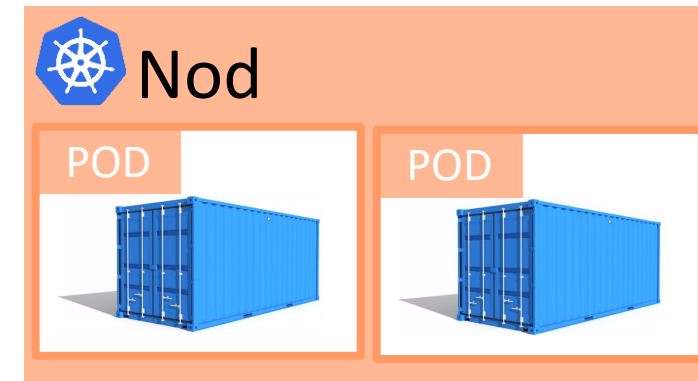
172.16.1.88 172.16.1.91



172.16.1.24 172.16.1.25

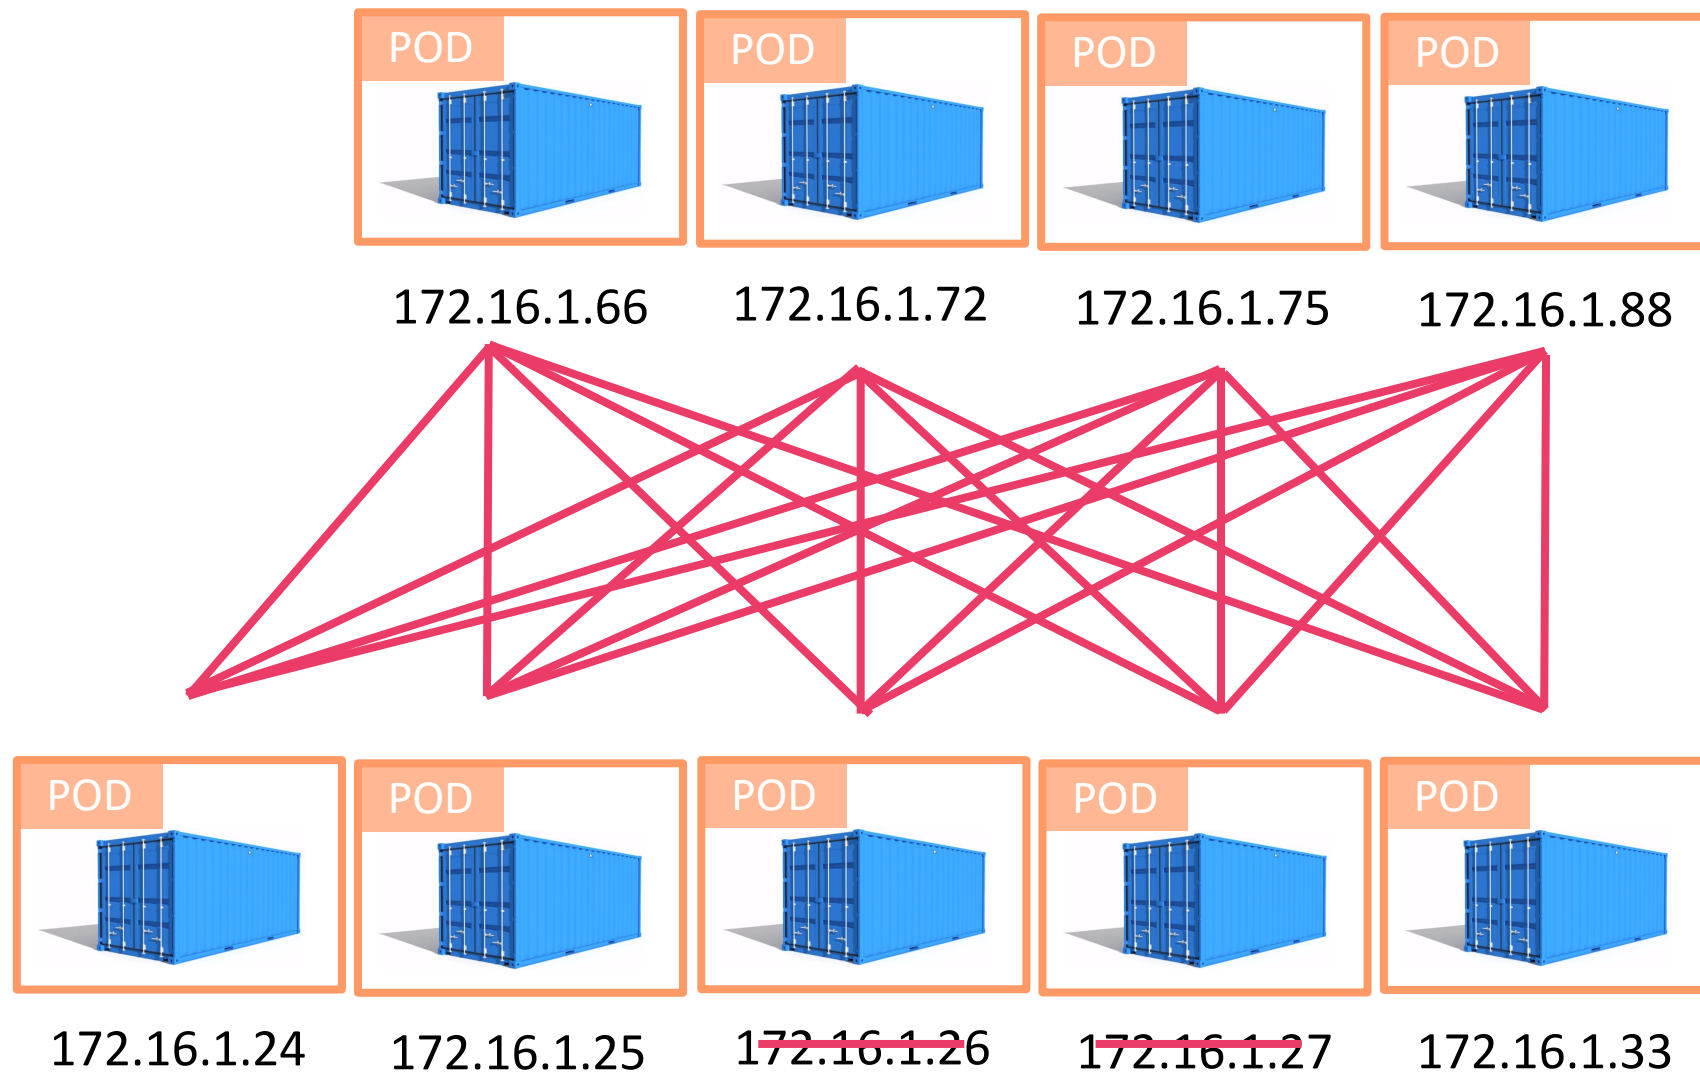


~~172.16.1.26 172.16.1.27~~

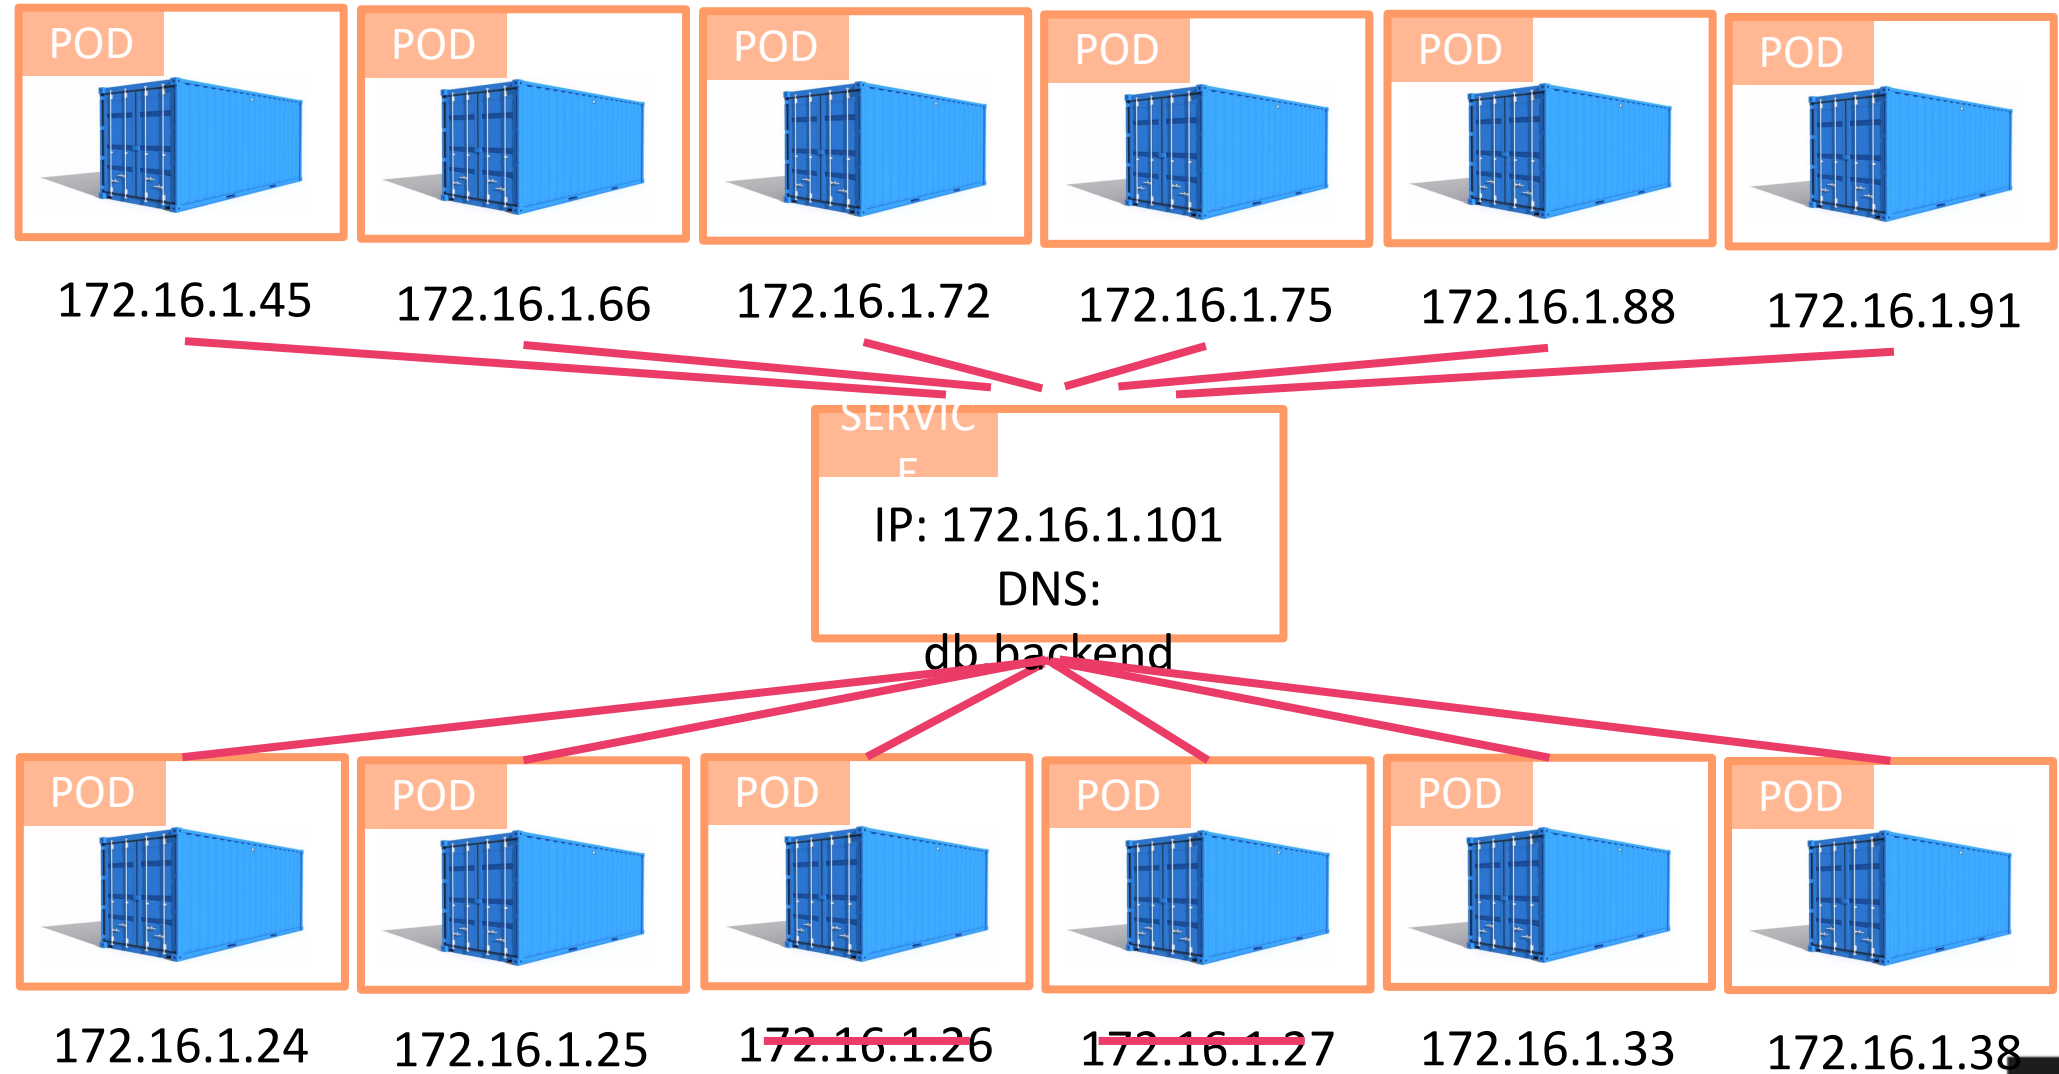


172.16.1.33 172.16.1.38

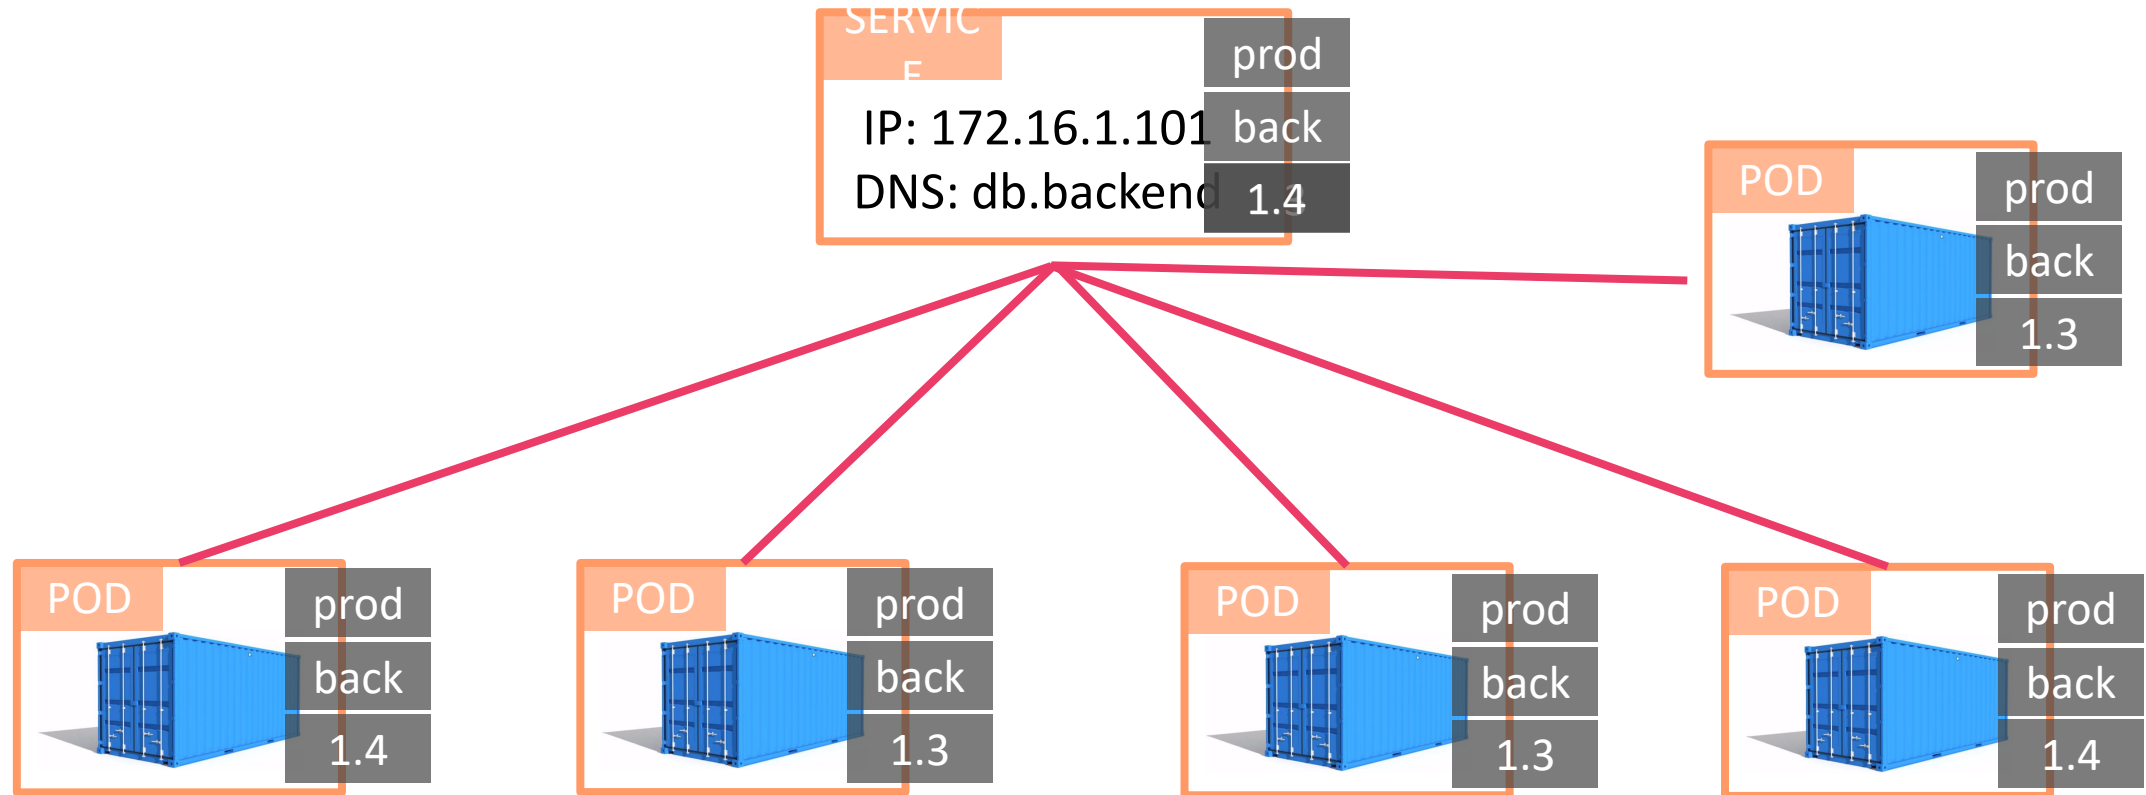
IP Churn



Services

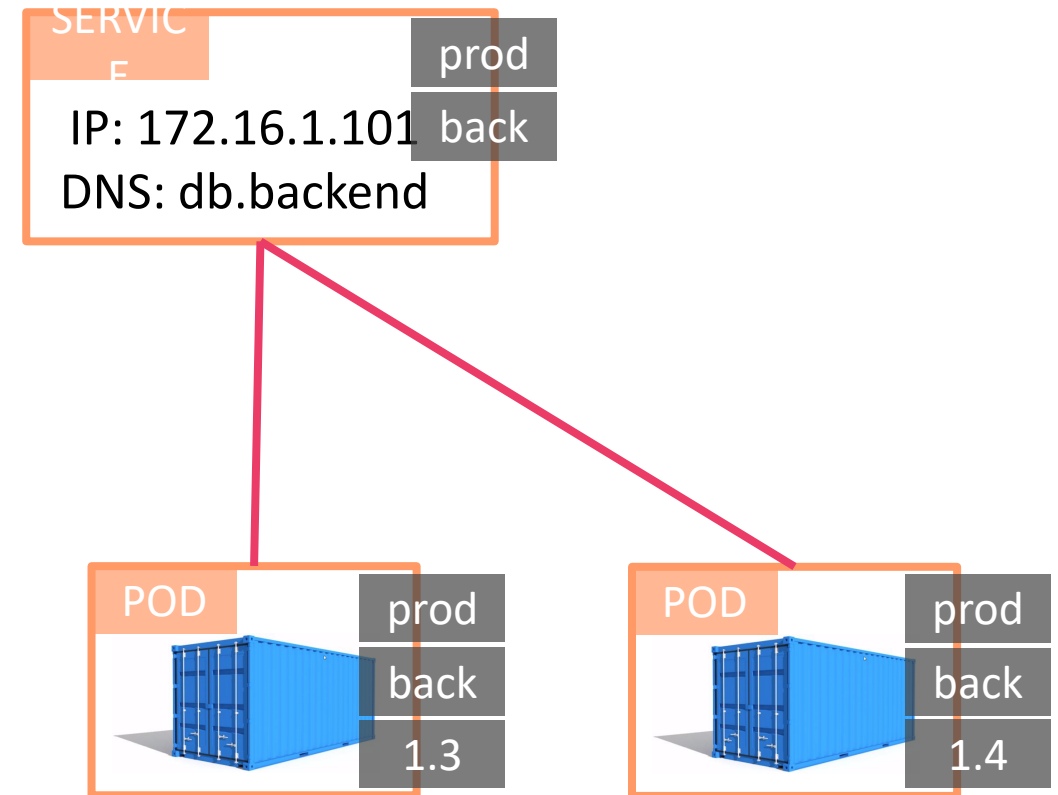


Labels



Services

- Only communicates to Healthy Pods
- Can be configured for Session Affinity
- Can point to resources outside the service
- Random Load Balancing
- Uses TCP by Default



nic

Resources

Slides and demos from the conference will be available at
github.com/nordicinfrastructureconference/2018 (bit.ly/2y7JhA3)